scoup.ai

scoup.ai marketing and go-to-market

2026-10-18

> **Pre-product. Synthetic data only. Not investment advice.**

Documents in this volume:

- marketing/marketing-plan.md
- marketing/gtm.md
- marketing/brand/README.md
- marketing/seo/keyword-research.md
- marketing/seo/content-plan.md
- marketing/seo/technical-checklist.md
- marketing/digital/channel-plan.md
- marketing/digital/cac-targets.md
- marketing/digital/creative-concepts.md

Generated from source by tools/pdf/build.py at source commit 3b18982d0662. Figures are drawn from source data at build time.

Contents

scoup.ai marketing plan (pre-product) 4
1. What we are marketing, and to whom 4
2. Positioning 5
3. Messaging by buyer 8
4. Launch plan: design-partner pilot first 9
5. Channel mix 9
6. Budget 10
7. How success is measured 12
8. Risks 12
9. Open questions 13

scoup.ai go-to-market (pre-product) 14
1. Principles 14
2. Segments 14
3. Winning in a contested segment 16
4. The first 100 customers, by profile 17
5. Design-partner recruitment plan 18
6. Qualification test 19
7. Segment summary 19
8. Open questions 20

scoup.ai brand summary (for the investor package) 21
1. The name 21
2. Logo and visual identity 21
3. Voice 22
4. Messaging pillars 22
5. Gaps 23

Keyword research: scoup.ai (B2B, US freight) 24
1. Method and its limits 24
2. What the Trends data says 24
3. Keyword clusters 26
4. What this means for the plan 30
5. How to replace UNKNOWN with real numbers 31
6. Open questions 31

Content plan: pillars, clusters and a 12-month calendar 33
1. What content has to do for us right now 33
2. Pillars and clusters 33
3. Twelve-month calendar (November 2026 to October 2027) 35
4. Which pieces win design partners, and how they are used 37
5. Original-data content from public sources (no customer data) 37
6. Rules every piece follows 39
7. How we will know whether it works 39
8. Open questions 39

Technical SEO checklist: audit of marketing/site/ 40
1. What was audited, and how 40
2. Findings summary 40
3. Findings and fixes 41
4. Launch gate 44
5. Open questions 44

Paid and performance channel plan for scoup.ai 45
Who we sell to first 45
Why paid channels come second 45
Channel 1: LinkedIn Ads (both segments; Year 1) 46
Channel 2: Google Search (both segments, high intent only; Year 1) 46

Channel 3: Industry events (Manifest in Year 1; TPM as the importer event) 47
Channel 4: Trade publications (Year 2) 48
Channel 5: Retargeting (both segments; Year 2) 48
Later phase: carriers and partner channels (not budgeted) 48
Summary and sequencing 49
Measurement 49
Open questions 49

CAC targets per channel for scoup.ai 50
Method 50
Funnel assumptions 50
CAC ceilings: what we can afford 50
Modeled CAC per channel 52
Targets we will hold the tests to 53
What would change these numbers 53
Open questions 53

Ad and creative concepts for scoup.ai 54
Audience and pillars 54
Rules every concept follows 54
Concept 1: "Off AP's desk" (lead test) 55
Concept 2: "The 30-day window" (lead test) 55
Concept 3: "No source, no dispute" (lead test) 55
Concept 4: "Fuel surcharge, recomputed" 56
Concept 5: "Too small for an audit program, too big to ignore" 56
Concept 6: "Import by container, ship by truck?" 56
Concept 7: "We're pre-product. Test us on your closed invoices." (founder content) 56
Concept 8: "The packet, unboxed" 57
Secondary creative 57
How we test them 57
Open questions 57

scoup.ai marketing plan (pre-product)

> **Status.** scoup.ai is pre-product: a prototype exists, it has run on synthetic data only, it has never been deployed, and nobody pays for it. This plan therefore markets a design-partner programme, not a product. It contains no traffic, conversion, recovery or revenue results, because none exist. Every target below is a pre-registered hypothesis, labelled as such. Name and voice follow `brand/README.md` (always lowercase **scoup.ai**). Budget and target CAC are set by `finance` on 2026-10-17 (section 6).
>
> Inputs: `freight-recovery/business/business-plan.md`, `business/investor-brief.md`, `marketing/seo-and-growth.md`, `marketing/landing/*`, `fundraising/README.md` (main checkout), the brand worktree (`freight-recovery-brand`, branch `feat/scoup-brand`), and competitor facts checked on vendors' own pages on 2026-10-17, plus `research/01-competitors.md`, `02-industries.md` and `04-gaps.md` (owner: `market-researcher`); where this file and those disagree, the research files win.

1. What we are marketing, and to whom

The product takes one load's rate confirmation, bill of lading or gate record, and invoice, applies the contract terms (clock start, free time, rounding, rate, cap, authorized accessorials), and produces an evidence packet plus a demand letter marked DRAFT that a person must approve before anything is sent. That is what the prototype does today on synthetic files (`freight-recovery-brand/marketing/site/pages/how-it-works.html`).

Run from the payer's side, the same engine checks a paid invoice: each accessorial and detention line against the rate confirmation that should have authorized it, and each container detention and demurrage (D&D) invoice against the FMC billing rule's content and timing requirements, with dispute deadlines tracked (`research/02-industries.md` notes the D&D checks and claims state machine exist on synthetic data).

Owner decision, 2026-10-17: shipper- and importer-first. The primary buyers are mid-market shippers (ASSUMPTION: $5M-$50M annual freight spend, `business-plan.md` section 2) that pay carrier and broker invoices carrying accessorial lines nobody authorized, and small and mid-size ocean importers paying D&D invoices they cannot check. The census counts 234,023 SME importers across all modes in 2024 (`research/02-industries.md`). `market-researcher` models 19,932 mid-market shippers from Census SUSB 2022 revenue classes (manufacturing 5,971, wholesale 8,798, retail 5,163; 9,366 to 35,349 depending on the freight-to-sales ratio assumed), and estimates the FMC D&D pool supports about 12,833 importers at the business case's size (`research/05-market-size.md`). Carriers, whom the business plan had proposed as the opening wedge, become a later channel reached through factoring and broker partners (`gtm.md`, Segment C). The payer holds the invoices and owns the money found, so a free review needs no counterparty's cooperation; the price of the choice is entering the side of the market where funded competitors already sell (section 2.2; `gtm.md` section 3).

The scale of the problem has one credible third-party anchor. ATRI's September 2024 detention study puts the cost of driver detention to US trucking at $15.1B a year; its release says "94.5 percent of fleets charge detention fees" yet are paid on "fewer than 50 percent of those invoices". `market-researcher` re-verified the figure against the live source on 2026-10-17 (`research/04-gaps.md`, `research/05-market-size.md`). It is a cost of waiting, not a pool of recoverable dollars, and external copy must say so whenever it quotes the number.

Outside voices describe the gap in the words we want to answer. A 2016 DAT survey of 257 carriers found "Only three percent of carriers were paid on 90 percent or more of their detention claims", and that two-thirds of brokers "paid detention only when they could collect the fee from the shipper or consignee". The survey is ten years old; we quote it as the only per-carrier measure found, not as current fact. A driver on TheTruckersReport (August 2025) describes the habit our packet formalises: "I always note my arrival and

departure times on my freight bill as well as a note for the reason". Broker detention terms vary ("They all have different policy and procedures"), so a packet must cite each counterparty's own terms rather than a generic rule.

Sources (quotes as recorded in `research/04-gaps.md`): - Source: DAT Solutions, "63 percent of drivers are detained for more than 3 hours per stop", 14 July 2016, https://www.dat.com/company/news-events/news-releases/63-percent-of-drivers-are-detained-for-more-than-3-hours-per-stop-dat-solutions, accessed 2026-10-17 - Source: TheTruckersReport forum, detention-denial poll thread, page 3, https://www.thetruckersreport.com/truckingindustryforum/threads/whats-the-1-reason-your-detention-claims-get-denied-poll.2529156/page-3, accessed 2026-10-17 - Source: TheTruckersReport forum, "Broker doesn't want to pay / one of those days", page 2, https://www.thetruckersreport.com/truckingindustryforum/threads/broker-doesn%E2%80%99t-want-to-pay-one-of-those-days.2519612/page-2, accessed 2026-10-17 - ATRI: full source line in `research/04-gaps.md` (release of 10 September 2024).

On the importer side, the FMC reports that the nine largest ocean carriers "have collected roughly $15.4 billion in D&D charges between April 1, 2020, and March 31, 2025" (about $3.08B a year, our arithmetic: $15.4B / 5), and the Harbor Trucking Association reported members disputing "more than 3 million detention and demurrage charges in 2020". Under the FMC billing rule, failure to meet its invoice content and timing requirements "relieves the billed party's obligation to pay", and the billed party has 30 calendar days to request mitigation, refund or waiver. Only 46 C.F.R. 541.4 was vacated by the D.C. Circuit in September 2025; the rest remains in force.

- Source: Federal Maritime Commission, Detention and Demurrage, https://www.fmc.gov/detention-and-demurrage/, accessed 2026-10-17
- Source: Supply Chain Dive, Harbor Trucking Association survey coverage, https://www.supplychaindive.com/news/detention-demurrage-port-congestion-fees-Harbor-Trucking-Association-TradeLanes/594406/, accessed 2026-10-17
- Source: Thompson Hine, "FMC Issues Final Rule on Billing Requirements and Practices for Detention and Demurrage Invoices", https://www.thompsonhine.com/insights/fmc-issues-final-rule-on-billing-requirements-and-practices-for-detention-and-demurrage-invoices/, accessed 2026-10-17
- Source: Thompson Hine, "Court sets aside part of FMC's demurrage and detention billing rule", 26 September 2025, https://www.thompsonhine.com/insights/court-sets-aside-part-of-fmcs-demurrage-and-detention-billing-rule/, accessed 2026-10-17

2. Positioning

2.1 Positioning statement

For transportation, import and AP leads at mid-market shippers and small and mid-size importers, who pay accessorial, detention and container D&D charges they cannot check against the document that authorized them, scoup.ai checks each line against the rate confirmation, contract or FMC billing rule, shows the mismatch and the arithmetic, and drafts the short-pay or dispute for a person to approve. Enterprise freight audit and payment platforms reconcile invoices at volume for large shippers; their smaller-shipper offers are subscriptions, and the ocean D&D specialist serves the largest importers. We take only the lines in question, charge nothing until money is recovered or avoided during the pilot, and hold back any line we cannot support instead of disputing it.

The statement makes four choices on purpose. It speaks to the payer, following the owner's shipper- and importer-first decision of 2026-10-17. It names documents rather than "data", because an AP manager recognises a rate confirmation and does not recognise a platform. It promises a checked line, not a recovery, because we have measured no recovery. And it defines the competition by what each option does with a disputed line, which is where we differ.

Carriers remain part of the story. The same engine builds a carrier's detention packet, and carriers become a later channel through factoring and broker partners (`gtm.md`, Segment C). The carrier hero line "Get paid for detention you can prove" stays on the carrier pages for that phase.

2.2 Named alternatives and where we stand against each

Each row quotes the vendor's own page or trade coverage, accessed 2026-10-17. "Where they are stronger" comes first because buyers will check it, and conceding it costs us nothing we could honestly claim.

Alternative	What they say they do (own words)	Where they are stronger	Our opening against them
Cass Information Systems (enterprise freight audit and payment)	"$37 billion in freight spend processed and paid annually"; "35 million freight invoices processed and paid each year"; logos include PepsiCo, Unilever, Toyota	Scale, payment processing, six decades in the business, a payment network of "15,000+ freight carriers"	Built for enterprise shippers' payables at Fortune 500 scale, by its own logos. A $10M-spend shipper is below that profile; we offer it a review of disputed lines with no programme to adopt.
Intelligent Audit (freight audit and payment)	"#1 Freight Audit & Payment Company"; "20% of Fortune 50 companies are customers"; services include "Freight Audit & Recovery", "DeepDetectAI", "Polly Carrier Agent"; FAQ: pricing is "a subscription model based on annual shipment volume, modes, and # of carriers"; for shippers under $10M a year, Catalyst offers "a free trial for 90 days" (launched June 2025 for parcel spend)	Breadth (parcel, claims, sourcing, visibility), AI features in market, an entry offer for smaller shippers	It is moving down-market on the shipper side, so we do not claim it ignores smaller shippers. It sells subscriptions, not contingency, and Catalyst started with parcel. Our difference: truckload and LTL accessorial and detention disputes, a per-line packet a counterparty can check, and a pilot paid only on money recovered. Detail in `gtm.md` section 3.
Loop (AI freight audit and payments)	"The industry's first Logistics Data Platform"; "99% of freight and parcel audits" automated; logos include General Mills and Macy's	AI-native shipper-side automation at volume; $160M raised, including a $95M Series C in April 2026; No. 10 on the 2026 Inc. 5000	Shipper-side and enterprise-weighted by its own logos. Its claim is audit throughput; ours is the evidence behind one disputed charge. The most direct funded competitor in our primary segment, and buying mid-size books (Data2Logistics, StrategIQ). Detail in `gtm.md` section 3.
TMS detention modules (McLeod Software, TMW Systems)	McLeod: carriers can "automatically notify and bill excessive detention time to customers"; TMW: "Detention Tracking" that can "automatically bill for detention time"	Already inside the carrier's system of record; no new vendor	They compute and bill the time. They do not assemble the rate-con clause, gate evidence and calculation steps into a packet built for a dispute, and they reach only carriers running those systems.
Telematics add-ons (FleetOptos Detention Tracker on Geotab)	"automates detention billing for trucking fleets using Geotab telematics data"; geofence dwell, contract billing rules, QuickBooks export	Automatic timestamps from the truck, which is strong evidence we do not generate	Tied to one telematics platform and focused on billing the charge. Its timestamps are an input we would cite; a partnership is likelier than a fight.
Hey Bubba (AI dispatcher with detention management, small carriers)	Records dwell from ELD geofence crossings, captures detention terms at booking, logs the claim with timestamps and notifies the broker while the truck is still at the facility; connects to "100+ ELDs"	Same buyer as our carrier segment, with the clock captured automatically inside a dispatch product used daily	It captures the clock. Its page does not describe a cited packet with the rate-con clause, BOL times and calculation, dispute deadlines, or the shipper side. A carrier on it is an integration prospect.
BlueCargo (ocean D&D audit and disputes)	Invoice checks "against OSRA invoice requirements"; "Dispute resolution with evidence-backed documentation"; claims "1 in 5 Top US Shippers"; named customers include GE Appliances and Samsung; $11M raised (2023)	The strongest D&D specialist: terminal and rail data connections, large-importer references	Serves the largest importers at major ports. Our opening is SME importers and one workflow covering container D&D and the trucking accessorials on the same shipments, which is untested demand.

Alternative	What they say they do (own words)	Where they are stronger	Our opening against them
Status quo (spreadsheets, email threads, write-offs)	n/a	Free, familiar, no data sharing	The real competitor. We win only if a packet takes less effort than the claim is worth, which is exactly what the design-partner programme measures.

Sources: - Source: Cass Information Systems, Transportation expense management, https://www.cassinfo.com/transportation-expense-management, accessed 2026-10-17 - Source: Intelligent Audit, home page, https://www.intelligentaudit.com/, accessed 2026-10-17 - Source: Loop, home page, https://loop.com, accessed 2026-10-17 - Source: Geotab Marketplace, FleetOptos listing, https://marketplace.geotab.com/solutions/fleetoptos/, accessed 2026-10-17 - Source: Bulk Transporter, "Module bills for downtime" (McLeod Software), https://www.bulktransporter.com/archive/article/21644416/wwwmcleodsoftwarecommodule-bills-for-downtime, accessed 2026-10-17 - Source: Intelligent Audit, FAQ, https://www.intelligentaudit.com/faq, accessed 2026-10-17 (via `research/01-competitors.md`) - Source: Hey Bubba, Detention Management, https://bubba.ai/solutions/detention-management, accessed 2026-10-17 (via `research/01-competitors.md`) - Source: Business Wire, Loop Series B, https://www.businesswire.com/news/home/20231003388240/en/, accessed 2026-10-17; AlphaMaven, "Loop raises $95M Series C", https://alpha-maven.com/story/venture-capital/loop-raises-95m-series-c, accessed 2026-10-17 (via `research/01-competitors.md`) - Source: BlueCargo, home page, https://www.bluecargo.io/home-tests, accessed 2026-10-17 (via `research/01-competitors.md`) - Source: Trucking Info, "TMW adds detention tracking feature", https://www.truckinginfo.com/94877/tmw-adds-detention-tracking-feature, accessed 2026-10-17 (older trade coverage; current product scope not re-checked)

Trax Technologies, CTSI-Global and nVision Global are enterprise freight audit and payment providers in the same position as Cass for our purposes: shipper-side, no carrier offer. BlueCargo serves ocean demurrage and detention, outside our truckload wedge for now. Full rows: `research/01-competitors.md`. The research finding that shapes our pricing message: no competitor profiled publishes prices or a contingency rate, and none of the freight audit and payment firms sells to carriers. With the shipper- and importer-first decision, the rows that matter most are Loop, Intelligent Audit (including Catalyst) and BlueCargo; the battlecard for winning against them is `gtm.md` section 3. The carrier-side rows (TMS modules, FleetOptos, Hey Bubba) matter again when the partner channel opens in phase 4.

2.3 Category and the word we own

We place scoup.ai in **accessorial, detention and D&D recovery for shippers and importers**, not "freight audit". Freight audit is the incumbents' category, defined on their terms, and entering it invites a comparison on scale we lose. The word we want attached to the brand is **proof**; the logo, two halves meeting exactly, draws the same idea (`brand/README.md`, section 2). The shipper hero line is **"Take accessorial review off AP's desk. Every line checked against the rate confirmation, with the source attached."**

We chose a workload line over the overbilling line "Stop paying accessorials that were never authorized" (drafted in `freight-recovery/marketing/landing/outline.md`, `/accessorial-audit/`) because the only shipper survey found points that way. In RXO's survey of 750 shippers and carriers, "61% of shippers agree that managing accessorial requests is one of the biggest time-wasters for their accounting team" and "71% are actively working on reducing accessorial spend", but "70% of shippers agree that carriers are usually honest when requesting accessorials — only 8% thought there was outright dishonesty" (`research/04-gaps.md`, gap 5). A line that implies carriers overbill argues with what most buyers believe; a line that removes work agrees with it, and the cited packet still shows any line that is not supported. RXO is a freight broker, not an audit vendor, so it has no audit product to sell, but it is not neutral either. The overbilling line stays as the A/B alternate in paid tests and as the H1 option for the `/accessorial-audit/` page, so the evidence can settle it.

Source: RXO, "What Shippers & Carriers Really Think About Trucking Accessorials (Original Research)", 25 November 2024, https://rxo.com/resources/research/freight-accessorials/, accessed 2026-10-17 The importer hero line is **"Check every D&D invoice before the 30-day window closes."** It names the deadline the FMC rule creates without claiming any charge is invalid; the supporting line under it is "We check each container detention and demurrage invoice against the FMC billing rule's content and timing requirements, and list what to dispute and by when." Both lines carry the source and the "confirm with counsel or your customs broker" note wherever the rule is described.

`seo` reports that buyers search in operational vocabulary ("rate confirmation", "detention pay") more than "freight audit", and that "freight claim" is read as cargo loss and damage (Google Trends, relative interest only, pulled 2026-10-17; data and script in `fr-wt/seo/marketing/seo/data/`). Headlines therefore use "accessorial", "detention", "rate confirmation" and "D&D invoice", and avoid "claim".

3. Messaging by buyer

The three pillars from `brand/README.md` section 4 (Proof attached; A person decides; Built for the loads big audit programs skip) stay constant. "Proof attached" leads for every buyer. "A person decides" answers trust objections; `market-researcher` found no outside evidence that buyers ask for approval controls (`research/04-gaps.md`), so it never leads.

Buyer (title)	What keeps them up	Core message	Proof we offer today	Objection we expect, and our answer
Shipper transportation / logistics manager (champion)	Accessorial requests eating the team's week; lines nobody can match to a rate con; carrier relationships	"Accessorial review off your team's desk: every accessorial and detention line checked against the rate confirmation, with the source attached and the few mismatches shown."	Synthetic example: linehaul above rate con, unauthorized lumper flagged for review, not auto-disputed	"We already have an audit vendor." Keep it. Send us only the lines you are unsure about.
Shipper AP manager / controller (user)	Accessorial requests as a time sink (61% of shippers in RXO's survey); paying what cannot be verified; audit trail at year-end	"Approve the supported lines in minutes, short-pay the rest with a packet that shows why, and keep the record for the auditors."	Confirmed vs "needs review" split; DRAFT letters; per-document fingerprints	"This will upset our carriers." Ambiguous lines go to a person, never disputed automatically.
Shipper CFO / VP finance (economic buyer)	Freight cost line; vendor risk; another subscription	"A review of three months of your paid invoices costs nothing. During the pilot, we are paid only on money recovered."	Pilot terms; the honest status page	"You have no customers." Correct, which is why the first step costs nothing and uses invoices you already have.
Importer import / logistics manager (champion)	D&D invoices arriving late or vague; dispute windows missed	"Check every container D&D invoice against the FMC billing rule's content and timing requirements before the 30-day window closes."	Synthetic D&D invoice check with the rule's required fields marked and the deadline shown	"Our forwarder handles this." Many forwarders pass charges through; we give you the list to dispute with them or through them.
Importer owner / CFO (economic buyer, smaller firms)	Cash leaking to port fees they cannot challenge	"Find the D&D charges the rule says you may not owe, and the truck accessorials on the same shipments, in one review."	The same synthetic samples; contingency terms	"Is this legal advice?" No. We check invoices against published requirements and draft; you or your counsel decide.
Factoring or broker partner (later phase, carrier channel)	Unpaid accessorial lines on purchased or brokered loads	"Your carriers' detention lines arrive with a packet you can forward to the shipper unchanged."	Sample carrier packet; written data-handling terms	"Will this create disputes?" Uncertain lines are held back, not claimed.

Every claim about the FMC rule in external copy carries the Thompson Hine source line from `research/04-gaps.md` (gap 3) and the note that 46 C.F.R. 541.4 was vacated in September 2025 while the content and dispute provisions remain in force. We describe what the rule requires; we never tell a buyer it does not owe a specific

charge, which is a legal judgement.

Messages we never use, for any buyer: a recovery percentage, a dollar figure from someone else's data, "AI-powered", "automated recovery", integration names, certification language, competitor claims beyond what their own pages state, or the unverifiable "~$3.1B freight audit and payment market" figure. The site's banned-phrase test (`freight-recovery-brand/marketing/site/test.mjs`) already enforces part of this list.

4. Launch plan: design-partner pilot first

A company with no product in market cannot "launch" in the usual sense, and pretending otherwise would burn the one asset we have, which is credibility with sceptical finance buyers. The launch is a sequence of gates taken from `business-plan.md` section 6; each phase spends only after the previous gate has evidence behind it. The audience in phases 1-3 is shippers and importers. Carriers enter in phase 4 through partners.

Phase	Gate it serves	Marketing goal	What we do	What we do not do	Exit evidence
0. Ready to talk (now)	Before G1	Be checkable when a prospect searches for us	Register the scoup.ai domain (brand Gap 10); set a real contact address (Gap 8); align site palette with brand (Gap 2); publish synthetic sample packets for a shipper accessorial review and a D&D invoice check; draft NDA and data-handling terms with counsel; answer a standard security questionnaire honestly	Paid acquisition, press, comparison pages	A finance buyer can read the site, the status page and a sample packet, and reach a person
1. Design-partner recruitment	G1 (leakage proven)	3-5 partners (at least two shippers, one importer) sharing three months of paid invoices	Founder-led warm intros via accountants, fractional CFOs, forwarders and customs brokers; individually written outreach (`gtm.md` section 5)	Bulk email, bought lists, logos	Signed NDAs, security questions cleared, invoices received
2. Invoice reviews	G1	A written review per partner	Run the prototype; hand back packets and, for importers, the FMC-rule issue list with deadlines; record where rules were wrong	Publish any partner number without written consent	Per-partner review; recoverable share measured, or the kill trigger hit
3. Contingency pilots	G3 (paid pilots)	3+ shipper or importer partners on live invoices	Pilot agreement at the contingency rate; human approval on every short-pay or dispute; first conversations with factoring and broker partners, to learn	Claim win rates before they are measured	Dollars recovered or avoided, cycle time, review minutes per invoice
4. First public proof and the carrier channel	After G3	One consented case study; the first partner channel to carriers	Case study; FMC-rule and accessorial content and calculator (`seo`); paid message tests; a co-offer with one factoring or broker partner	Scale paid spend before CAC is measured	First inbound leads attributable to content; one signed partner

Phases 0-2 depend on the security and legal items in the brand worktree's `security.html` "Not yet" list, and shippers will ask about them sooner than carriers would. No real document is accepted before handling terms are signed, and the status page says real customer data is not accepted until those items are done. Marketing cannot shorten that dependency; it can only avoid promising around it.

5. Channel mix

Founder time is the scarce input (`seo-and-growth.md` section 11), so the mix is ranked by expected design-partner conversations per founder hour, not by reach. The shares are planned allocations of founder marketing time in phases 1-3, a decision rather than a measurement.

Channel	Role	Planned share of founder marketing time	Starts	Notes
Warm introductions and individually researched outreach	Primary source of shipper and importer design partners	45%	Phase 1	Accountants, fractional CFOs and personal network first; named-account list from the profiles in `gtm.md` section 4; CAN-SPAM and GDPR/ePrivacy rules apply
Referral partners: customs brokers and freight forwarders (importers); later factoring firms and brokers (carriers)	Introductions to many accounts through one trusted party	20%	Phase 1 for forwarders; phase 3-4 for carrier partners	Never imply a partnership before one is signed
Founder LinkedIn (posts and substantive comments)	Credibility with finance and logistics titles; inbound	15%	Phase 1	Synthetic accessorial and D&D checks worked through step by step
Owned content and SEO (owner: `seo`)	Discovery for accessorial-dispute and FMC D&D queries	10%	Phase 0 foundations, phase 4 cadence	Detailed plan in `marketing/seo/**`
Events and associations	Meeting supply-chain, import and finance buyers in person	5%	Phase 1-3	Attend before sponsoring; one trade show pass in year 1
Paid search and LinkedIn ads (owner: `digital-marketing`)	Message tests only	5%	Phase 4 (after G3)	Targeting shipper and importer titles; budget in section 6

The order matters more than the percentages. Introductions and referral partners come first because they reach people who already hold paid invoices; content comes second because it makes those conversations easier; paid comes last because paying for clicks before we know which message converts would measure nothing.

5.1 Lead magnets for the new audience

The calculator spec in the main checkout (`freight-recovery/marketing/landing/calculator-spec.md`) was written carrier-first; its shipper variant has only two inputs and there is no importer variant. We drafted a fuller shipper variant (accessorial and detention lines on paid invoices) and a new importer D&D variant (invoices at stake and the share that would pass the 30-day window unreviewed) in `brand/calculator-variants.md`, with worked examples and acceptance tests. The owner of `marketing/landing/` should merge them before the calculator is built.

6. Budget

Year-1 marketing programmes: $30.0K ($7.5K per quarter). Sales travel: $60.0K, a separate line. Total Year-1 sales and marketing: $90.0K. No marketing hire until year 3.

These are `finance`'s plan inputs for a pre-product company, not measurements (`fr-wt/finance/financial/budget.md`; `model.xlsx`, Assumptions `o_mkt` and Unit Economics tab). Legal work (NDA, data-handling terms, pilot agreement, claims review) sits in G&A under a $150K legal and regulatory line, not in marketing. Marketing programmes by year in the base case ($K): 30 / 90 / 200 / 320 / 420 / 480 / 520.

Within the $30.0K, we allocate as follows. `finance` confirmed this split on 2026-10-17, on the condition that the $7,500 of paid tests stays locked until gate G3.

Line	Phase	What it buys	Amount
Domain registration, site fixes, sample packet page	0	scoup.ai domain (brand Gap 10), brand-consistent site, contact address, synthetic example	$2,000

Line	Phase	What it buys	Amount
Domain-expert reviewer for explainers and packet format	0-2	A freight-billing practitioner checks our rules and copy	$4,000
CRM, outreach tooling and list data	1	Tracking every conversation and its source; shipper and importer account list	$3,000
Content production and design (packet template, deck template; brand Gap 9)	0, 4	Branded artifacts a buyer actually handles	$6,000
One trade show pass (Manifest, attendance, not a booth)	1-3	In-person meetings with carriers, factors and associations	$1,995
Paid message tests: LinkedIn $4,500 + Google $3,000	4	Search and LinkedIn ads on a handful of buyer-intent phrases	$7,500
Reserve	all	Unplanned content or event costs	$5,505
Total			**$30,000**

Arithmetic: 2,000 + 4,000 + 3,000 + 6,000 + 1,995 + 7,500 + 5,505 = $30,000. Founder travel to design partners and the assumed $1,500 travel to the trade show come from the separate $60.0K sales line, which `finance` scopes to "travel to US design partners and freight-industry meetings".

After the owner's shipper- and importer-first decision, `digital-marketing` keeps year 1 at $9,495 (LinkedIn $4,500 split 60/40 shipper/importer, Google $3,000 in one shipper and one importer ad group, the $1,995 Manifest pass) and cuts year 2 from $5,500 to $4,000: one trade newsletter $3,000 (JOC if importers lead, DC Velocity if shippers do) and retargeting $1,000. The $1,500 Meta test for carriers is dropped entirely, and the carrier creative is withdrawn. Paid tests in both years stay locked until G3. On events, we follow `digital-marketing`'s recommendation: TPM 2027 (JOC, Long Beach, 28 February-3 March 2027; fee unknown) replaces Manifest only if importers produce more design-partner conversations than shippers by January 2027; we do not add a second event. Both figures are inside `finance`'s programmes line (year 2: $90K).

Earlier, `digital-marketing` proposed $16,495 for paid tests and events (`fr-wt/digital/marketing/digital/channel-plan.md`). Its LinkedIn ($4,500), Google ($3,000) and Manifest pass ($1,995; its $1,500 travel estimate moves to the sales-travel line) fit inside the $30.0K. Its remaining lines were re-planned after the owner's decision, as set out in the paragraph above. If phase 4 slips past year 1, the $7,500 moves into year 2's $90K.

Target CAC (sales and marketing cost divided by new accounts won, base case): $13.4K in year 2, $15.0K in year 3, $17.9K in year 5. In year 3 that equals $0.68 of sales and marketing per $1 of new ARR, against a Benchmarkit 2024 median of $2.00. At the median, CAC would be $43.8K and the base case would need a $12.3M seed (`finance`). The plan therefore depends on three cheap acquisition paths staying cheap: founder-led selling, the free invoice review, and referral partners (forwarders and customs brokers now; factoring firms and brokers for the carrier channel later). Those three are exactly what phases 1-3 test, and the CAC actually observed on the first 20 accounts replaces this target.

Source: Benchmarkit, 2025 B2B SaaS benchmarks, https://www.benchmarkit.ai/2025benchmarks, accessed 2026-10-17 (via `fr-wt/finance/financial/unit-economics.md`)

`finance` also gives a sourced pessimistic CAC of $31K per customer (media cost only, excluding sales labour), about twice the plan's ~$15K. Paid results in phase 4 are therefore read as a message test, not as a growth channel to scale.

Source: Metadata.io, 2026 B2B Paid Media Benchmark Report (2025 data, 153 advertisers), https://metadata.io/md/benchmark-report-2026.md, accessed 2026-10-17 (via `fr-wt/finance/financial/assumptions.md`, input `pess_cac`)

The business plan's earlier illustrative CAC of ~$25K and LTV:CAC of ~5.2x (`investor-brief.md` section 6.5) are superseded by `finance`'s figures above.

Creative concepts are mapped to the three pillars in `fr-wt/digital/marketing/digital/creative-concepts.md`.

Spending rule regardless of the total: nothing in phase 4 is released until G3 evidence exists, and legal lines are paid before any acquisition line.

7. How success is measured

At this stage the metric that matters is qualified conversations that turn into shared files, not traffic. Targets below are pre-registered hypotheses for phases 1-3, set before any data so that we cannot move them afterwards. If they are missed, the plan is revised, not the record.

Stage	Metric	Pre-registered target (hypothesis)	Source of truth
Reach	Named accounts contacted with a personalised message	60 accounts in phase 1	CRM
Response	First calls held	15 calls (one in four contacted)	CRM
Qualification	Calls meeting the qualification test in `gtm.md` section 6	8 qualified	CRM notes
Commitment	Signed NDA and paid invoices received	3-5 design partners (at least two shippers, one importer)	Signed agreements
Value	Invoice reviews delivered with partner feedback recorded	One per partner	Review documents
Conversion	Partners moving to a contingency pilot	3 pilots (G3)	Pilot agreements
Learning	Rule corrections logged from partner feedback	Every correction logged; no target count	Issue tracker
Channel attribution	Source of each partner conversation	100% of conversations tagged	CRM field

The funnel ratios in that table are our own planning guesses, not benchmarks: 60 contacted, one in four to a call (15), about half of calls qualifying (8), and three to five of those signing. Once phase 1 ends, the observed ratios replace them and set the phase 3 plan. Web metrics (indexed pages, calculator completions) belong to `seo` and `digital-marketing` and are reported, not targeted, until two months of real data exist.

8. Risks

1. **Credibility gap.** Polished marketing for an unvalidated product erodes trust. We keep the pre-product notice on every page and lead with the sample packet, not claims.
2. **Contested segment.** The owner's shipper- and importer-first decision moves us onto ground where Loop ($160M raised), Intelligent Audit (with Catalyst for smaller shippers) and BlueCargo (ocean D&D) already sell. We compete only where the buyer wants no platform fee, no integration and per-line evidence, and offer to coexist with an existing audit vendor (`gtm.md` section 3). If that niche is too small, phase 1 will show it through the stop rule.
3. **Shippers fear their carriers.** A shipper may hesitate to short-pay carriers it depends on. Human approval and the "needs review" category are the answer.
4. **Longer, heavier sales.** Shippers bring a finance approver and security questionnaires; the honest status page will be read closely. The recruitment timeline in `gtm.md` section 5 allows for it.
5. **Nobody shares files.** If design partners will not share paid invoices under NDA, G1 fails. The kill trigger in `business-plan.md` section 6 applies; marketing does not paper over it.
6. **Recoverable share is small.** The plan's economics rest on an unmeasured recovery rate. Phase 2 exists to measure it before any spend scales. For importers, `research/02-industries.md` puts the SME D&D ceiling at $0.93M a year by year 5 on assumed inputs, so D&D is sold as the importer's way in, not as a business of its own.

7. **Competitor drift.** Intelligent Audit is moving down-market; Loop is buying mid-size books (Data2Logistics, StrategIQ); Hey Bubba already sells detention capture to carriers we will later reach through partners. Our differentiation is narrow and must be re-checked every quarter.
8. **Trademark and domain.** No trademark search for "scoup.ai" is recorded (`brand/README.md`, Gap 7); it must close before printed material or paid media. The scoup.ai domain also appears unregistered (RDAP 404 on 2026-10-17, reported by `seo`; brand Gap 10); registering it is a phase 0 item.

9. Open questions

- The freight-to-sales ratio behind the 19,932 shipper count (from 2008), and the share of SME importers that import by ocean container.
- Whether shippers accept contingency against Loop's platform fees and Intelligent Audit's subscriptions, and whether the coexist offer is enough to start.
- Whether importers want container D&D and truck accessorials in one workflow (untested).
- Whether customs brokers and forwarders will refer importers, and on what terms.
- For the later carrier channel: factoring and broker partner terms, and whether Hey Bubba or FleetOptos would partner (their timestamps in, our packets out) rather than compete.
- Pricing: the financial model (`fr-wt/finance/financial/assumptions.md`) sets a 25% pilot contingency for shippers and importers, then an annual platform fee of $6K rising to $17K plus 17% falling to 12%, and 20% for carriers. All are ASSUMED; design-partner conversations test them, and any change goes through `finance` so the package keeps one price path.
- Who serves as the freight-domain reviewer for copy and rules.

scoup.ai go-to-market (pre-product)

> **Status.** Pre-product: synthetic data only, never deployed, no customers, no revenue. Everything below is a plan. Sales cycles, conversion rates and price acceptance are hypotheses to be measured in the design-partner programme, and each is labelled. Segment counts come from `research/02-industries.md` (owner: `market-researcher`); only counts that drive a GTM decision are repeated here, with their source lines. Positioning and messaging: `marketing-plan.md` sections 2-3. Name: always lowercase **scoup.ai** (`brand/README.md`).
>
> **Owner decision, 2026-10-17: shipper- and importer-first.** Mid-market shippers and small and mid-size ocean importers are the primary segments. Carriers move to a later phase, reached through factoring and broker partners. This reverses the business plan's carrier-first working hypothesis (`business-plan.md` section 2), and this file is rewritten accordingly.

1. Principles

Three rules shape every segment. First, we sell a free review of past invoices before we sell anything else, because a vendor with no customers earns trust only by doing work on the buyer's own files. Second, pricing is contingency on dollars actually recovered or avoided, so the buyer's downside is the time it takes to share closed invoices (`business-plan.md` section 4). Third, the founder sells until roughly the first 20 accounts (`business-plan.md` section 5); no sales hire before CAC is measured.

Why shipper- and importer-first suits that model: the buyer holds the invoices and pays them, so a review needs no counterparty's cooperation, and the money found is the buyer's own spend. The cost of the choice is competition. Unlike the carrier side, where no freight audit and payment firm sells (`research/01-competitors.md`), the shipper side is contested by Loop and Intelligent Audit, and ocean detention and demurrage (D&D) by BlueCargo. Section 3 sets out how we compete there.

2. Segments

Segment A: mid-market shippers (primary)

- **Profile.** US manufacturers and distributors with $5M-$50M annual freight spend (ASSUMPTION band from `business-plan.md`), mostly truckload and LTL, paying carrier and broker invoices through AP either without a freight audit programme or with an audit vendor that does not argue disputed accessorial lines.
- **Size.** 19,932 firms (manufacturing 5,971; wholesale 8,798; retail 5,163), average freight spend about $9.53M. This is a modelled count: `market-researcher` converted the $5-50M freight-spend band to a $100M-$1B revenue band at an assumed 4.68% freight-to-sales ratio and counted Census SUSB 2022 firms in those classes. At 2% the count is 9,366; at 10% it is 35,349 (`research/02-industries.md`, `research/05-market-size.md`). There is no independent shipper survey on accessorial disputes. Source: U.S. Census Bureau, Statistics of U.S. Businesses 2022, https://www2.census.gov/programs-surveys/susb/tables/2022/us_naicssector_large_rcptsize_2022.xlsx, accessed 2026-10-17
- **Buyer titles.** Director or manager of transportation/logistics (champion); AP manager or controller (daily user and approver); CFO or VP finance (economic buyer).
- **Trigger events.** A freight-cost overrun flagged at month-end close; a year-end audit questioning accessorial spend; a carrier or broker contract renegotiation; dropping or switching a freight audit vendor; a new controller asking why accessorials are paid without support.
- **Channel.** Founder LinkedIn aimed at logistics and finance titles; warm introductions through accountants, fractional CFOs and supply-chain associations; search content on accessorial disputes (owner: `seo`); later, introductions from carrier partners on shared lanes.

- **Sales motion.** Free review of three months of paid carrier invoices against the rate confirmations that authorized them; confirmed findings separated from items needing review; then a contingency pilot on new invoices. Expect a second approver (finance) and a security questionnaire before any live data.
- **Price point.** Pilot: 25% of dollars recovered or short-paid with support, no platform fee. From model year 2: hybrid, an annual platform fee of $6K per account rising to $17K by year 7, plus a contingency share falling from 17% to 12%. All ASSUMED in the financial model (`financial/assumptions.md`, inputs `share` and `fee`), which governs every price in this package. No freight competitor publishes a contingency rate (`research/01-competitors.md`), so there is no market check yet.
- **Sales cycle (hypothesis).** 8-16 weeks from first call to signed pilot: two to four weeks for NDA and security questions, three to six for the invoice review, two to four for pilot terms with finance. Our estimate, not a benchmark.

Segment B: small and mid-size ocean importers (primary, paired with A)

- **Profile.** US manufacturers and wholesalers with fewer than 500 employees that import in ocean containers and receive container detention and demurrage invoices from ocean carriers and terminals. Many also ship domestically, so they belong to Segment A as well; that overlap is the point.
- **Size.** 234,023 SME importers in 2024 across all modes (Census); the number importing by ocean container is UNKNOWN — needs primary research. `market-researcher` estimates that the FMC D&D pool supports only about 12,833 importers at the business case's exposure per account (`research/05-market-size.md` section 5), which is the more useful planning number. Source: U.S. Census Bureau, "A Profile of U.S. Importing and Exporting Companies, 2023-2024", release CB26-73, https://www.census.gov/foreign-trade/Press-Release/edb/edbrel2024.pdf, accessed 2026-10-17
- **Buyer titles.** Import or logistics manager (champion); AP manager or controller (user); owner or CFO (economic buyer at the smaller end).
- **Trigger events.** A D&D invoice spike after port congestion; an invoice arriving late or missing required fields; a 30-day dispute window about to close; a forwarder passing D&D charges through without detail.
- **Why the trigger is real.** Under the FMC billing rule, "Failure of the billing party to abide by the invoice content and timing requirements relieves the billed party's obligation to pay", and the billed party "has 30 calendar days from the date of issuance to request mitigation, refund, or waiver". After the D.C. Circuit's September 2025 ruling, only 46 C.F.R. 541.4 was vacated; the rest remains in force. Source: Thompson Hine, "FMC Issues Final Rule on Billing Requirements and Practices for Detention and Demurrage Invoices", https://www.thompsonhine.com/insights/fmc-issues-final-rule-on-billing-requirements-and-practices-for-detention-and-demurrage-invoices/, accessed 2026-10-17 Source: Thompson Hine, "Court sets aside part of FMC's demurrage and detention billing rule", 26 September 2025, https://www.thompsonhine.com/insights/court-sets-aside-part-of-fmcs-demurrage-and-detention-billing-rule/, accessed 2026-10-17
- **Channel.** Customs brokers and freight forwarders as referral partners. They hold the importer relationship and the documents, and `market-researcher` found none selling priced D&D recovery. They are also a substitute, and when a forwarder acts as an NVOCC it is itself the billing party, so a referral from it is unlikely for its own invoices (`research/01-competitors.md`). We approach customs brokers first and forwarders only for charges billed by ocean carriers and terminals; importer associations and trade events; search content on the FMC rule and the 30-day window; founder LinkedIn.
- **Sales motion.** Free check of the last 90 days of D&D invoices against the rule's content and timing requirements, with a dispute deadline list; then a contingency pilot. Nothing is disputed without the importer's approval.
- **Price point.** Same path as Segment A: 25% of charges waived or refunded in the pilot, then the hybrid platform fee plus 17% falling to 12% (ASSUMED, `financial/assumptions.md`).
- **Sales cycle (hypothesis).** 4-10 weeks: the 30-day window creates urgency, but the smallest importers rely on their forwarder and may need its cooperation. Our estimate.

- **Economics caution.** `research/02-industries.md` estimates the SME D&D slice alone at a $0.93M a year ceiling by year 5 on assumed inputs, and concludes D&D is "a feature that deepens an importer account, not a standalone business". We sell it as the importer's way in, alongside domestic accessorial review, not as a separate product line.

Segment C (later phase): carriers, through factoring and broker partners

- **Profile.** US for-hire truckload carriers with 5-100 power units: 72,053 active interstate carriers filed an MCS-150 update since 17 October 2024 (FMCSA census, queried by `market-researcher`). Source: FMCSA, Company Census File (dataset az4n-8mr2), https://data.transportation.gov/resource/az4n-8mr2.json, accessed 2026-10-17
- **Why later, and why through partners.** The owner chose to sell first to the party that pays. Carriers remain valuable: the same engine builds their detention packets, and no freight audit and payment firm serves them. But a founder selling one carrier at a time cannot reach 72,053 accounts. Factoring firms, which buy carriers' invoices and absorb unpaid accessorials, and brokers, two-thirds of whom "paid detention only when they could collect the fee from the shipper or consignee" (DAT, 2016), can each bring many carriers at once. Source: DAT Solutions, 14 July 2016, https://www.dat.com/company/news-events/news-releases/63-percent-of-drivers-are-detained-for-more-than-3-hours-per-stop-dat-solutions, accessed 2026-10-17
- **Partner titles.** Factoring: head of operations or carrier relations; collections manager. Brokers: VP operations; carrier-relations or accessorials team lead. 28,438 active brokerage authorities in January 2024 (`research/02-industries.md`). Source: FreightCaviar, "Brush Pass Research: 7.8% fewer brokerages than a year ago", 24 January 2024, https://freightcaviar.com/brush-pass-research-7-8-fewer-brokerages-than-a-year-ago, accessed 2026-10-17
- **Motion.** After shipper pilots exist, offer partners a co-offer: their carriers get packets a broker can forward to the shipper unchanged. Referral share or platform fee agreed only after a partner pilot; no exclusivity before G3.
- **Price point.** Carriers: 20% contingency, no platform fee (ASSUMED, `financial/assumptions.md`, input `car_share`). Partner terms: UNKNOWN — needs primary research in partner conversations.
- **Timing.** Partner conversations may start in phase 3 to learn; signed partnerships target phase 4 or later.
- **Caution.** DAT's 2016 survey found only 20% of responding brokers ranked detention a top-five problem, against 84% of carriers. Brokers may feel this less than the carriers they pay.

3. Winning in a contested segment

The shipper side is where funded competitors already sell. We do not try to out-feature them. We win deals where four conditions hold, and we walk away where they do not: the buyer wants no platform fee until money is found; the buyer cannot or will not integrate a TMS or ERP feed; the pain is disputed truckload or LTL accessorials and D&D lines rather than parcel or payment processing; and the buyer wants to see, per line, why a charge is wrong before short-paying it.

Competitor	Their offer to our segment (sourced)	Where they win	Where we win	Our talk track
Loop	AI freight audit and payments; "99% of freight and parcel audits" automated; $160M raised; model described by Sacra as "platform fees with usage- or volume-based components" (analyst inference); merged with Data2Logistics (founded 1964, Fort Myers, FL; mid-market book) on 18 August 2025 and acquired StrategIQ; says implementation takes "as fast as 31 days" (`research/01-competitors.md`)	Buyers who want audit, payment and visibility in one platform; buyers with data feeds ready to connect	No platform fee before money is found; no integration for the pilot (we work from the invoices and rate confirmations the buyer already has); a per-line evidence packet for each disputed charge	"Keep your audit and pay process. Send us the lines you are unsure about, and we will show you, line by line, which ones the rate confirmation does not support."

Competitor	Their offer to our segment (sourced)	Where they win	Where we win	Our talk track
Intelligent Audit Catalyst	For companies shipping under $10M a year, "a free trial for 90 days" then "one of our three Catalyst plans" (prices not published); launched June 2025; the Catalyst page says "Simple. Quick. Affordable." and lists "Automated parcel audit with direct carrier integrations" first, with nothing on truckload accessorials or D&D; main offer priced as "a subscription model based on annual shipment volume, modes, and # of carriers"	Parcel-heavy shippers; buyers who want breadth (claims, sourcing, visibility); a known brand	Truckload and LTL accessorial and detention disputes; a contingency pilot instead of a subscription trial; shippers above the $10M Catalyst line who find an enterprise programme too heavy	"If parcel is your main spend, Catalyst may suit you better. If your problem is truckload accessorials and detention you cannot verify, that is all we do."
BlueCargo (ocean D&D)	Invoice checks "against OSRA invoice requirements", "Dispute resolution with evidence-backed documentation"; claims "1 in 5 Top US Shippers"; named customers include GE Appliances and Samsung; "95% Dispute Success" (vendor claim, grade C)	Large importers at major ports; terminal and rail data connections	SME importers below its customer profile; one workflow covering container D&D and the truck detention and accessorials on the same shipments (whether importers want them combined is untested)	"Check every D&D invoice against the FMC rule before the 30-day window closes, and handle the trucking side of the same shipments in the same packet."

Two adjacent players matter on the importer side. Portcast offers a D&D invoice module and a free carrier-by-carrier D&D rate tool, a likely lead magnet for the same SME importers; our importer calculator (`brand/calculator-variants.md`) answers a different question, what is at stake and what would pass the 30-day window. Forwarders and customs brokers run dispute intake but sell no priced D&D recovery that `market-researcher` found; they are both a channel and a substitute (Segment B).

Sources: `research/01-competitors.md` and `research/04-gaps.md`, which carry the full source lines, including: - Source: Intelligent Audit, FAQ, https://www.intelligentaudit.com/faq, accessed 2026-10-17 - Source: Sacra, Loop company profile, https://sacra.com/c/loop, accessed 2026-10-17 - Source: BlueCargo, home page, https://www.bluecargo.io/home-tests, accessed 2026-10-17 - Source: FreightWaves, "Loop and Data2Logistics merge to unlock profit in the supply chain", https://www.freightwaves.com/news/loop-and-data2logistics-merge-to-unlock-profit-in-the-supply-chain, accessed 2026-10-17 - Source: Portcast, home page, https://www.portcast.io/, accessed 2026-10-17

Rules for using this table: competitor facts are re-checked on the vendor's own page the week they are used with a prospect; we never claim a competitor's performance beyond what it publishes; we say where they are stronger before we say where we differ; and we never imply a competitor's customers are unhappy. The "coexist" message matters most: a shipper with an audit vendor already can add us for disputed lines only, so our first deal does not require displacing anyone.

4. The first 100 customers, by profile

No company is named as a target. Naming real shippers or importers before they agree to talk would be a list of strangers, and the package's synthetic-data rule rules out picking them from real shipment records. In the business plan's base case the company ends Year 3 with 80 paying accounts and Year 4 with 180 (`investor-brief.md` section 9.2, illustrative), so account 100 falls early in Year 4 of that model.

Cohort	Accounts	Mix (our planning split)	Typical profile	How they arrive	Who sells
1. Design partners turned pilots	1-5	3 shippers, 2 importers	Shippers with $10-30M freight spend whose AP team has disputed accessorials before; importers receiving D&D invoices at least monthly, ideally also shipping domestically by truck	Warm introductions through accountants, fractional CFOs, forwarders and personal network	Founder

Cohort	Accounts	Mix (our planning split)	Typical profile	How they arrive	Who sells
2. Founder-led early accounts	6-20	9 shippers, 6 importers	Same profiles, widening to $5-10M shippers whose controller handles freight invoices directly	Referrals from cohort 1, customs-broker and forwarder introductions, founder LinkedIn, association events	Founder, plus the first customer-success hire when review volume demands it
3. Repeatable accounts	21-100	43 shippers, 22 importers, 15 carriers	Shippers and importers reached through search content and referral partners; carriers brought in batches by one factoring or broker partner	Content, referral partners, small paid tests after G3, the later carrier partner channel	First dedicated sales hire (planned 9-18 months post-seed, `investor-brief.md` section 8.3) plus founder

Arithmetic: shippers 3 + 9 + 43 = 55; importers 2 + 6 + 22 = 30; carriers 0 + 0 + 15 = 15; total 100. An importer that also ships domestically counts once, under the pain that brought it in. The split is our planning choice following the owner's decision; it is not a forecast.

Common traits across the shipper and importer accounts, used as the qualification test in section 6: they pay carrier, broker or ocean-carrier invoices themselves; they can export three months of paid invoices with the rate confirmations or contracts behind them; a named person will approve every short-pay or dispute; and they accept contingency pricing.

5. Design-partner recruitment plan

Goal. 3-5 design partners, at least two shippers and at least one importer, each sharing three months of paid invoices (or 20-50 closed loads) read-only under an NDA agreed first.

What partners get. A review of their own invoices: each accessorial, detention and D&D line checked against what authorized it, confirmed findings separated from items needing review, the arithmetic shown, and for importers a list of FMC-rule issues and open dispute deadlines. Nothing to buy. What we get: corrections to our rules, the measured recoverable share for gate G1, and permission (if granted) to use anonymised results later.

Preconditions before the first outreach message. 1. The scoup.ai domain registered and a real contact address on the site (brand Gaps 8 and 10). 2. NDA, data-handling, retention and deletion terms drafted by counsel (legal sits in G&A, `finance`). 3. A synthetic sample packet for a shipper accessorial review and one for a D&D invoice check. 4. A recorded decision with counsel on how partner files are handled before the security status page's "Not yet" items close.

Sourcing, in priority order. 1. Personal and second-degree network: anyone running transportation, AP or imports at a manufacturer or distributor. 2. Accountants and fractional CFOs who serve mid-market manufacturers; one introduction each to a client with visible freight cost. 3. Two or three customs brokers or forwarders, asked to introduce importers who complain about D&D. 4. Founder LinkedIn posts walking through a synthetic accessorial check and a synthetic D&D invoice check against the FMC rule. 5. Individually researched outreach to companies matching the cohort 1 profile, through their published business contact. No purchased lists, no automated sequences.

Outreach sequence (per account). Message 1 names the likely situation (accessorials paid without support, or D&D invoices close to the dispute window), asks for a 20-minute call, states the pre-product status in the first two lines, and asks them not to send documents yet. One follow-up after a week with the matching synthetic sample. Then stop. Opt-outs recorded (CAN-SPAM; GDPR/ePrivacy for EU/EEA contacts, since the operator is in Norway).

The first call (agenda). How invoices are approved today; how many carrier, broker and D&D invoices a month; where the rate confirmations and contracts live; what happened to the last disputed line; whether an audit vendor is already in place (if so, the coexist offer); whether they would share three months under an NDA. The sample packet comes last.

Timeline (hypothesis, from the start of phase 1).

Weeks	Milestone
1-2	Preconditions done; target list of 60 profile-matched shippers and importers
3-7	Outreach; 15 first calls held
6-10	NDAs and security questions cleared with 3-5 partners; invoices received
9-14	Reviews delivered; partner feedback logged; G1 decision
14-18	Pilot terms offered where findings are material

The timeline is two weeks longer than a carrier-first plan would need, because shippers bring a finance approver and a security questionnaire. The 60 / 15 / 3-5 figures are the pre-registered targets in `marketing-plan.md` section 7.

Stop rule. If 60 personalised contacts produce fewer than two NDAs, we stop and revisit the offer and the segment before contacting more accounts.

6. Qualification test

A shipper or importer qualifies for a design partnership or pilot only if all five hold. The list follows the business plan's qualification (`investor-brief.md` section 6.3), adapted to the payer side.

1. Pays carrier, broker or ocean-carrier invoices itself, with rate confirmations, contracts or tariffs in writing.
2. Willing to share three months of paid invoices (or 20-50 closed loads) read-only under an NDA.
3. Plausibly material leakage: accessorials or detention paid without checking the authorizing document, or D&D invoices paid without checking the FMC rule's requirements.
4. A named person who approves every short-pay or dispute.
5. Accepts that the product is pre-product, the review runs on our prototype, and nothing is sent on its behalf.

7. Segment summary

	A. Mid-market shippers (primary)	B. SME ocean importers (primary)	C. Carriers via partners (later)
Size	19,932 (modelled; 9,366-35,349 by freight-to-sales ratio)	234,023 SME importers, all modes; ~12,833 supported by the FMC D&D pool	72,053 carriers with 5-100 trucks; 28,438 brokerages
Economic buyer	CFO / VP finance	Owner or CFO	Factoring or broker head of operations (partner)
Champion / user	Transportation manager; AP manager	Import/logistics manager; AP	Collections or carrier-relations lead
Trigger	Freight overrun at close; audit questions; vendor change	D&D spike; non-compliant invoice; 30-day window	Rising accessorial disputes on purchased or brokered loads
Channel	LinkedIn, accountants and fractional CFOs, associations, search	Customs brokers and forwarders, importer events, search	Partner introductions, after shipper pilots exist
Motion	Free 3-month invoice review, contingency pilot, then hybrid	Free 90-day D&D check, contingency pilot	Co-offer after a partner pilot
Price (ASSUMED, financial model)	25% pilot; from year 2, $6K-17K annual platform fee + 17% falling to 12%	Same as A	20% for carriers, no platform fee; partner terms UNKNOWN

	A. Mid-market shippers (primary)	B. SME ocean importers (primary)	C. Carriers via partners (later)
Sales cycle (hypothesis)	8-16 weeks to pilot	4-10 weeks to pilot	3-6 months to partner agreement
Main competitors	Loop; Intelligent Audit (incl. Catalyst)	BlueCargo; Intelligent Audit D&D validation	Hey Bubba, FleetOptos, TMS detention modules

Price source: `fr-wt/finance/financial/assumptions.md` (inputs `share`, `fee`, `car_share`), which supersedes the earlier `business-plan.md` section 4 path (20%, then $1-2K a month plus 10-12%). The only public contingency reference finance found is in parcel audit (Refund Retriever, 40-50%), a different market. Contingency is an established model in the industry; one audit firm states its "compensation is limited to a contingent share of the actual savings generated by our services". Source: Lojistic, Freight audit FAQ, https://www.lojistic.com/shipping-faqs/ftl-ltl-ocean-rail-air-freight-audit-questions, accessed 2026-10-17

8. Open questions

- The shipper count (19,932) rests on a 4.68% freight-to-sales ratio from 2008; refreshing that ratio is the first sizing task (`research/02-industries.md`). No independent shipper survey on accessorial disputes exists, so design-partner calls are the first demand evidence on this side.
- Share of SME importers that import by ocean container: UNKNOWN — needs primary research.
- Whether importers want container D&D and truck accessorials handled in one workflow (untested; `research/04-gaps.md`).
- Whether shippers accept a 25% pilot contingency against Loop's platform fees and Intelligent Audit's subscriptions, and whether a coexist offer (disputed lines only) is enough to start.
- Whether customs brokers and forwarders will refer importers, and on what terms.
- Factoring and broker partner terms for the later carrier channel.
- Target CAC by segment. `finance` sets a blended target ($13.4K in year 2, $15.0K in year 3; `marketing-plan.md` section 6); pilots should record cost per account by segment.

scoup.ai brand summary (for the investor package)

> Status: summary of decisions already made in the brand worktree (`C:\Users\mariu\ventures\freight-recovery-brand`, branch `feat/scoup-brand`, head `3933aae`). This file does not redesign anything. Where the brand work and this package disagree, the brand worktree wins and the disagreement is listed under "Gaps". The company is pre-product: no hosted service, no customers, synthetic data only.

1. The name

The brand worktree settles the name as **scoup.ai**, written in lowercase every time, and treats the company and the product as one brand (`brand/README.md`, "Name rules"). The rules we carry into every marketing document:

- Write **scoup.ai**. Never "Scoup", "Scoup.ai", "SCOUP.AI" or "Scoup AI". Possessive: **scoup.ai's**.
- Avoid opening a sentence with the name; rephrase. If it cannot be avoided, keep it lowercase.
- Never place the name inside an all-caps context (CSS `text-transform: uppercase` included).
- In page and app headers the logo replaces the text name; the text name stays in titles, alt text, footers and legal text.
- "freight invoice recovery", "freight audit", "detention recovery" are ordinary lowercase words, not the name.
- Code and repositories keep the technical name (`freight-recovery`, `@fr/*`, `freight_recovery`). Only the name people see changed (commit `98c9a54`, "Rebrand to scoup.ai").

The integration owner confirmed lowercase **scoup.ai** package-wide on 2026-10-17 (`BRIEF.md`, line 13).

2. Logo and visual identity

The logo is the **"Reconciled s"**: an s built from two halves that meet exactly at the spine, standing for the invoice and the contract agreeing. The upper half is neutral ink, the lower half teal, and the s sits on a rounded tile (corner radius 22% of its width) that doubles as app icon, avatar and favicon (`brand/README.md`, opening paragraph and "Do and don't"). The idea matters for marketing because it is the product's claim drawn as a shape: two documents checked against each other until they agree.

Assets (all in `freight-recovery-brand/brand/`): horizontal and stacked lockups for light and dark surfaces, one-colour black and white lockups for print, wordmark, mark (24 px and up) and a heavier `favicon.svg` for 16-24 px. Clear space is 1x the wordmark's x-height around any lockup and 0.25x the tile width around the tile alone. Minimum horizontal lockup: 20 px tall on screen, 6 mm in print.

Palette, taken from existing product tokens, with contrast ratios as measured in the brand README:

Role	Light surfaces	Dark surfaces
Ink (letters, tile)	`#0f172a` (17.85:1 on white)	`#e2e8f0` (15.19:1 on `#0b1220`)
Teal (lower half of the s, single accent)	`#0f766e` (5.47:1 on white)	`#2dd4bf` (10.06:1 on `#0b1220`)
Muted (".ai")	`#64748b` (4.76:1 on white)	`#8b9ab0` (6.55:1 on `#0b1220`)

The web app carries the same system in `web/src/styles/tokens.css`: an 8 px spacing grid, slate neutrals, one accent (deep teal `#0f766e`), semantic colours used sparingly, 120-180 ms ease-out motion, Inter Variable as the UI typeface, and WCAG AA contrast enforced by a unit test (`web/src/styles/contrast.test.ts`).

Prohibited uses (from "Do and don't"): offsetting or recolouring the two halves, the dark tile on dark slate (1.05:1), the s without its tile in colour, stretching, shadows, gradients, animation, and setting the name in capitals or another typeface as a stand-in for the logo.

3. Voice

The brand worktree has no written voice guide, but the shipped marketing site (`freight-recovery-brand/marketing/site/pages/*.html`) shows a consistent voice that we treat as decided.

1. **Plain and specific.** Sentences name documents a freight person handles every day: "rate confirmation, bill of lading and invoice". Arithmetic is shown, not summarised ("Free time 2 h, rate $60/h, cap $480 … 3 h 10 min beyond free time", `how-it-works.html`).
2. **Candid about status.** Every page carries "pre-product, in design-partner validation", and the security page opens with "We would rather tell you what is missing than imply a posture we do not have" (`security.html`).
3. **Limits stated as promises.** "We never send anything without human approval"; "We do not move money, and we do not give legal advice" (`index.html`, "What we do not do").
4. **Asking, not selling.** The call to action is a design-partner conversation where "honest feedback is the point" (`index.html`, `contact.html`).
5. **No proof we do not have.** No logos, testimonials, savings percentages, "AI-powered" or certification claims; the site's `test.mjs` fails the build on a banned-phrase list (`marketing/site/README.md`).

Lines in use: the hero **"Get paid for detention you can prove."** (`index.html`); for carriers, "Bill earned detention with the proof attached instead of writing it off"; for shippers, "Catch unauthorized or unsupported accessorials; ambiguous lines are flagged for human review, not auto-disputed". The earlier business plan proposed "The proof standard for freight recovery" (`freight-recovery/business/investor-brief.md`, section 6.1); the site never used it.

After the owner's shipper- and importer-first decision (2026-10-17), the primary hero lines are the shipper line "Take accessorial review off AP's desk. Every line checked against the rate confirmation, with the source attached." (chosen 2026-10-17 over "Stop paying accessorials that were never authorized", which stays as the A/B alternate; reasoning in `../marketing-plan.md` section 2.3) and a new importer line, "Check every D&D invoice before the 30-day window closes" (`../marketing-plan.md`, section 2.3). The carrier hero stays for the later partner phase.

4. Messaging pillars

The site copy and the investor brief's section 6.1 converge on three pillars. We adopt them in the wording below; `../marketing-plan.md` section 3 maps them to each buyer.

Pillar	One line	What we can show today	What we may not say
Proof attached	Every claim carries the documents, the clause and the arithmetic behind it.	Evidence packet with per-document SHA-256 fingerprints and printed calculation steps, on synthetic data (`how-it-works.html`).	Any recovery rate, win rate or dollar figure.
A person decides	Software drafts; a person approves every letter, and uncertain lines are held back, not claimed.	Letters generated as "DRAFT / not sent"; findings split into confirmed and "needs human review" (`index.html`).	"Automated recovery", "hands-free".
Built for the loads big audit programs skip	Mid-market shippers and small and mid-size importers who cannot justify an enterprise audit program; carriers later, through partners (owner decision, 2026-10-17).	Design-partner offer on closed loads, nothing to buy (`contact.html`).	That incumbents cannot serve them, or any competitor claim without a dated source.

Pillar two is a trust answer, not a lead. `market-researcher` found no outside evidence that buyers ask for approval controls (`research/04-gaps.md`), so "Proof attached" leads every message and "A person decides" answers the objections that follow.

The investor brief's pillar "AI you can audit" and its neural-mesh speed figures do not carry into marketing. Those figures come from internal benchmarks on synthetic workloads and tell a buyer nothing they can check; under the package's honesty rule they stay out of external copy.

5. Gaps

1. **Name casing (resolved 2026-10-17).** `BRIEF.md` first wrote "Scoup"; the integration owner confirmed lowercase **scoup.ai** package-wide and updated the brief. Older repo documents still say "Freight Recovery" and need the same pass before external use.
2. **Marketing site palette does not match the brand.** `marketing/site/static/style.css` still uses a blue accent (`--accent:#0b5cad`) and a system font stack, while the brand and app use teal `#0f766e` and Inter. The site was rebranded in name and logo only.
3. **No written voice guide.** Section 3 is reconstructed from shipped copy. A one-page guide with do/don't sentence pairs would prevent drift once more people write copy.
4. **No messaging for factoring companies, brokers or 3PLs**, although the business plan names them as channels.
5. **Print sizes for the stacked lockup and the mark** are "not yet specified" in `brand/README.md`.
6. **Strokes not outlined.** The brand README's production TODO requires converting strokes to outlines before files go to a printer or sign maker.
7. **Trademark clearance for "scoup.ai"** is not recorded. UNKNOWN — needs primary research (USPTO, EUIPO and Norwegian IPO searches, ideally by counsel).
8. **Contact address unset.** `contact.html` renders `__CONTACT_EMAIL__` with an `example.invalid` fallback, so outreach cannot point at the site yet.
9. **Only a social card exists** (`static/og-card.png`). There is no deck template, letterhead, or branded template for the evidence packet itself, which is the artifact buyers will see most.
10. **Domain not registered.** On 2026-10-17 `seo` found that scoup.ai does not resolve and the .ai registry RDAP lookup returned 404 "Object not found"; the site currently lives at socrates6969.github.io/freight-recovery/. The name and the domain must be secured together before any copy, card or ad prints the domain.
11. **"Claim" is a loaded word in search.** `seo` reports that "freight claim" means cargo loss and damage claims (parcel carriers dominate it), and that carrier vocabulary ("detention pay", "rate confirmation") draws more search interest than "freight audit" (Google Trends, relative interest only; script in `fr-wt/seo/marketing/seo/data/`). Page titles and headlines should say "detention", "accessorial" and "rate confirmation" rather than "claim"; body copy may still use "claim" where the meaning is clear.

Keyword research: scoup.ai (B2B, US freight)

> **Status.** scoup.ai is pre-product: synthetic data only, never sold, no customers, no live product. This file builds on `marketing/seo-and-growth.md` section 3 and replaces its `TBD: measure` placeholders where a citable source now exists. It contains **no absolute monthly search volumes and no keyword-difficulty scores**, because we have no paid SEO tool account (Semrush, Ahrefs, Moz or similar) and no Google Ads account with spend, which is what unlocks exact Keyword Planner figures. Every absolute volume is therefore `UNKNOWN — needs primary research`. What we do have is **Google Trends relative interest**, pulled on 2026-10-17 with a reproducible script, which tells us how terms compare with each other but not how many searches any of them gets.
>
> **Owner decision, 2026-10-17:** scoup.ai goes to market shipper- and importer-first; carriers become a later or partner channel. The clusters below are ordered by that decision, not by measured search interest. Section 4 explains where the two diverge.

1. Method and its limits

We pulled Google Trends data for 29 candidate terms in seven batches of five (batches F and G were added after the owner's shipper- and importer-first decision, same day, same settings) (US, 2021-10-01 to 2026-09-30, all categories, web search). Every batch carries the same anchor term, "freight audit", so a term's interest can be expressed as a multiple of the anchor and compared across batches. The script and its raw output sit next to this file: `marketing/seo/data/trends_pull.py` writes one CSV per batch plus `trends_summary.csv`, and `marketing/seo/data/trends_related.py` writes the related-query lists to `trends_related.csv`. Re-running them will give slightly different numbers, because Trends samples its data; the order of magnitude should hold.

Google's own description of the data sets the limits of what we can claim. Each data point is divided by total searches for the geography and period, the result is scaled 0 to 100, and "search terms with low volume appear as '0'". A Trends value is a share, not a count. A ratio of 2.33 means a term drew about 2.3 times the search interest of "freight audit" over the five years; it does not mean either term gets any particular number of searches.

Source: Google, FAQ about Google Trends data, https://support.google.com/trends/answer/4365533, accessed 2026-10-17 Source: Google Trends (via the `pytrends` client), US, 2021-10-01 to 2026-09-30, pulled 2026-10-17; batch A as an explore URL: https://trends.google.com/trends/explore?date=2021-10-01%202026-09-30&geo=US&q=freight%20audit,detention%20pay,detention%20time,freight%20claim,accessorial%20charges, accessed 2026-10-17

Difficulty is harder still. Without a tool we cannot score it, so the difficulty column below is a qualitative judgement from sampling who appears in web-search results for the query on 2026-10-17. Those samples come from a general web-search API, not a captured Google results page, and they should be treated as a hint about the kind of site that competes (carrier factoring firms, TMS vendors, enterprise audit firms, law firms), not as a ranking.

2. What the Trends data says

The data cuts against the owner's chosen segments on the software terms and is kinder on the problem terms. The vendor-category phrases a mid-market shipper would type when shopping ("freight audit software", "freight bill audit", "freight invoice audit") sit at or below the Trends reporting threshold, at 0 to 2 percent of the interest in plain "freight audit". The carrier's own vocabulary ("detention pay", "detention charges", "rate

confirmation", "TONU", "layover pay") draws between 0.7 and 3.3 times the interest of "freight audit". People search for the problem in the words of the person who does the waiting, and almost nobody searches for the software category.

The shipper and importer picture improves once we stop looking for software phrases. Single words a shipper's AP team uses on an invoice do draw interest: "freight bill" 3.88 times "freight audit", "freight invoice" 1.47 times, "accessorial" 0.88 times (against 0.08 for "accessorial charges"). On the importer side, "demurrage" is the second largest term we measured (4.88 times, broad), while the specific phrases are small: "demurrage charges" 0.22, "container detention" 0.13, "demurrage and detention" 0.10, "fmc demurrage" 0.03, "per diem charges" 0.02. The importer searches for the meaning of a charge and for a carrier's tariff far more than for the rule or for a tool.

Term	Batch	Mean interest in batch	Ratio to "freight audit"	Reading
freight audit (anchor)	all	4.4 to 8.5	1.00	Anchor; its batch value moves because Trends rescales each request
detention time	A	21.9	4.97	**Polluted**: includes school and legal "detention"; do not read as freight demand
demurrage	D	34.2	4.88	Broad ocean-shipping term; related queries are definitional ("demurrage meaning")
freight payment	B	39.2	4.63	Likely includes payment-portal navigation; vendor/brand intent mixed
rate confirmation	E	15.2	3.34	Carrier and broker document term; strong fit for evidence content
tonu	C	23.5	2.77	Ambiguous abbreviation; freight share unknown
detention pay	A	10.2	2.33	Carrier intent; top related query "detention pay trucking"
freight claim	A	6.6	1.51	**Wrong intent**: related queries are FedEx/UPS loss-and-damage claims
detention charges	C	10.5	1.24	Mixed carrier/shipper; closest to our product language
layover pay	C	6.0	0.71	Driver-pay intent; adjacent
detention fee	E	1.6	0.36	Small but on-topic
lumper fee	C	2.1	0.25	Small; maps to the product's "flag for human review" rule
demurrage and detention	D	0.7	0.10	Small; ocean/intermodal; FMC-rule audience
freight dispute	E	0.5	0.10	Small
accessorial charges	A	0.4	0.08	Near threshold
detention trucking	D	0.5	0.07	Near threshold (people say "detention pay trucking" instead)
freight audit software	B	0.2	0.02	Near threshold
freight bill audit	B	0.1	0.01	Near threshold
freight invoice audit	B	0.0	0.00	Below Trends threshold
freight bill	G	32.8	3.88	Shipper AP term; likely includes carrier bill-pay navigation
lumper	G	32.9	3.89	Ambiguous single word; freight share unknown
freight invoice	G	12.4	1.47	Shipper AP term
accessorial	G	7.4	0.88	People search the bare word, not "accessorial charges"
demurrage charges	F	1.8	0.22	Importer intent
container detention	F	1.1	0.13	Importer intent
fmc demurrage	F	0.3	0.03	Rule-aware importers; small
per diem charges	F	0.2	0.02	Near threshold
accessorial fees	D	0.0	0.00	Below Trends threshold

Term	Batch	Mean interest in batch	Ratio to "freight audit"	Reading
freight overcharge	E	0.0	0.00	Below Trends threshold

Source: Google Trends, seven batches as above, `marketing/seo/data/trends_summary.csv` (A-E) and `marketing/seo/data/trends_summary_F_importer_G_shipper.csv` (F-G), pulled 2026-10-17. Ratios are our arithmetic: term mean divided by the "freight audit" mean in the same batch (example, batch A: 10.2 / 4.4 = 2.33 for "detention pay"). Ratios built on small means are coarse; a 0.0 means "too little data for Trends to report", not "zero searches".

The related-query lists sharpen the picture. For "freight audit" the top related queries are "freight audit and payment" (100), "freight audit services" (91) and "freight audit software" (76), which tells us the buyer who does search this space is looking at outsourced audit-and-pay firms, the incumbents' home ground. For "detention pay" the top related query is "detention pay trucking" (100) followed by "detention pay for truck drivers" (33), a driver and small-carrier audience. "Demurrage" produced a long tail of definitional and carrier-specific queries ("demurrage meaning", "demurrage vs detention", "msc detention and demurrage tariff" rising +650 percent), a sign that ocean D&D searchers want explanations and tariff lookups. "Freight claim" is dominated by "fedex freight claim" (100) and "ups freight claim" (32), which is cargo loss-and-damage, a different job from ours.

Source: Google Trends related queries, US, 2021-10-01 to 2026-09-30, `marketing/seo/data/trends_related.csv`, pulled 2026-10-17. Values are Trends' relative scores (top: 0-100; rising: percent growth), not volumes.

3. Keyword clusters

Each cluster below names its intent, the evidence we have, a qualitative difficulty call with its reasoning, and the page it should map to (page plan in `content-plan.md`). Absolute volume and tool difficulty are blank for every row for the reason given at the top. Clusters are numbered in the owner's priority order (shipper, then importer, then the shared invoice and category clusters, then carriers). Each also fits something the prototype does today: compute detention from rate confirmation, gate times and invoice, flag unauthorized accessorials for human review, check ocean D&D invoices against 46 CFR 541 (build branch, synthetic data), and produce an evidence packet with a DRAFT demand. The cluster order is not a forecast of traffic.

The numbering changed on 2026-10-17: old Cluster 2 (accessorials) is now Cluster 1, old Cluster 5 (D&D) is now Cluster 2, and old Cluster 1 (carrier detention) is now Cluster 5.

Cluster 1: Accessorial and freight-invoice disputes (shipper side), priority 1

The shipper-side vocabulary is thin in Trends ("accessorial charges" 0.08x, "accessorial fees" below threshold), but the bare words an AP team uses are not ("accessorial" 0.88x, "freight invoice" 1.47x, "freight bill" 3.88x). This cluster leads because the owner chose shippers first, and it is about conversion quality more than reach: the few people who search "how to dispute accessorial charges" are exactly the mid-market AP or logistics manager we want as a design partner. Difficulty looks medium-low. The sampled results are guides from a 3PL and broker (Worldwide Express), a cross-border carrier (Nuvocargo), a managed-transportation firm (Translogistics) and a warehouse encyclopedia (Racklify). They give generic advice (dispute in writing, attach the BOL and POD); none walks through one invoice line against one rate-confirmation clause with the arithmetic shown.

Keyword	Intent	Trends ratio	Monthly volume	Difficulty (qualitative)	Target page
how to dispute accessorial charges / dispute accessorial fees	Informational, shipper	not pulled (parent 0.08)	UNKNOWN — needs primary research	Medium-low	`/guides/dispute-accessorial-charges/`
unauthorized accessorial charges / lumper fee dispute	Informational	lumper fee 0.25	UNKNOWN — needs primary research	Low	`/guides/lumper-fees/`

Keyword	Intent	Trends ratio	Monthly volume	Difficulty (qualitative)	Target page
accessorial audit / accessorial charge audit	Vendor-aware	not pulled	UNKNOWN — needs primary research	Low-medium	`/accessorial-audit/`
how to dispute a detention charge (shipper)	Informational	not pulled	UNKNOWN — needs primary research	Low-medium	`/guides/dispute-a-detention-charge/`
accessorial / accessorial meaning	Informational	0.88	UNKNOWN — needs primary research	Medium: carriers' own accessorial schedules likely rank	`/guides/accessorials/` (hub)
freight invoice / freight bill (how to read and check one)	Informational, shipper AP	1.47 / 3.88	UNKNOWN — needs primary research	Medium; bill-pay navigation dilutes intent	`/guides/how-to-check-a-freight-invoice/`

Source: Worldwide Express, How to dispute accessorial fees, https://www.wwex.com/shipping-resources/how-to-dispute-accessorial-fees, accessed 2026-10-17 Source: Nuvocargo, Accessorial charges in freight (2026 guide), https://nuvocargo.com/blog-posts/accessorial-charges-freight-cost, accessed 2026-10-17 Source: Translogistics, Common carrier disputes and how to solve them, https://www.translogisticsinc.com/blog/carrier-disputes, accessed 2026-10-17 Source: Racklify, How to audit and dispute a carrier invoice, https://racklify.com/encyclopedia/how-to-audit-and-dispute-a-carrier-invoice-practical-steps-for-warehouses/, accessed 2026-10-17

Cluster 2: Demurrage and detention for importers, priority 1

D&D is a real and regulated pain, and the TypeScript build (branch `feat/hardening`, read-only) already contains an ocean D&D invoice-audit module behind the feature flag `intelligence.ocean_dd`, run on synthetic data only. Its rules quote 46 CFR 541.5 to 541.8 verbatim and check invoice contents and deadlines field by field (`packages/shared/src/ocean-dd/regulation.ts`). An earlier draft of this file said the prototype did not compute D&D; that was wrong. The Federal Maritime Commission's billing rule (46 CFR part 541, published 26 February 2024, effective 28 May 2024) sets minimum invoice contents, a 30-day invoicing deadline and at least 30 days for the billed party to dispute. In September 2025 the D.C. Circuit vacated only 46 CFR 541.4 (who may be billed); the content and dispute provisions stay in force. The owner chose importers as a first segment, and research gives the reasons to care: 234,023 small and medium US importing companies in 2024 (all modes; the ocean-container share is unknown), and an FMC statement that the nine largest ocean carriers collected about $15.4 billion in D&D charges from April 2020 to March 2025. Searchers want definitions ("demurrage meaning", "demurrage vs detention") and carrier tariffs, so the cluster earns an explainer and a checklist that mirror what the module checks. Difficulty is medium: law firms (HFW, Squire Patton Boggs, Thompson Hine), trade associations and ocean-visibility vendors (BlueCargo, Portcast) publish here. Pages may describe the module only as pre-product software tested on synthetic invoices, with no accuracy or recovery claim.

Keyword	Intent	Trends ratio	Monthly volume	Difficulty (qualitative)	Target page
demurrage vs detention / demurrage meaning	Informational	demurrage 4.88 (broad)	UNKNOWN — needs primary research	Medium	`/guides/demurrage-vs-detention/`
demurrage and detention / D&D charges	Informational	0.10	UNKNOWN — needs primary research	Medium	same page
demurrage charges / demurrage fees	Informational to commercial	0.22	UNKNOWN — needs primary research	Medium	`/guides/demurrage-charges/`
container detention / per diem charges	Informational	0.13 / 0.02	UNKNOWN — needs primary research	Medium	`/guides/demurrage-vs-detention/`

Keyword	Intent	Trends ratio	Monthly volume	Difficulty (qualitative)	Target page
[carrier] detention and demurrage tariff (e.g. "msc detention and demurrage tariff", rising)	Navigational, tariff lookup	related query only	UNKNOWN — needs primary research	Low for a neutral index; carriers rank for their own	`/guides/dd-free-time-tracker/`
how to dispute demurrage charges / D&D refund request	Commercial, importer	not pulled	UNKNOWN — needs primary research	Medium	`/guides/dispute-demurrage-charges/`
FMC demurrage detention rule / D&D invoice requirements	Informational, compliance	not pulled	UNKNOWN — needs primary research	Medium: law firms	`/guides/fmc-dd-invoice-checklist/`

Source: HFW, FMC's new demurrage and detention regulations become effective May 2024, https://www.hfw.com/insights/fmcs-new-demurrage-and-detention-regulations-become-effective-may-2024, accessed 2026-10-17 Source: Squire Patton Boggs, Federal Maritime Commission issues final demurrage and detention rule, https://www.squirepattonboggs.com/en/insights/publications/2024/03/federal-maritime-commission-issues-final-demurrage-and-detention-rule, accessed 2026-10-17

Source: Thompson Hine, Court sets aside part of FMC's demurrage and detention billing rule (26 September 2025), https://www.thompsonhine.com/insights/court-sets-aside-part-of-fmcs-demurrage-and-detention-billing-rule/, accessed 2026-10-17 (via `research/02-industries.md` and `research/04-gaps.md`)

The eCFR page redirected us to a bot check on 2026-10-17. The build module's source comment says 541.4 was removed from the CFR effective 29 December 2025 by Federal Register document 2025-23920; we have not read that document ourselves, so it must be checked before a D&D page cites the removal date.

Source: U.S. Census Bureau, A profile of U.S. importing and exporting companies, 2023-2024 (CB26-73, 5 May 2026), https://www.census.gov/foreign-trade/Press-Release/edb/edbrel2024.pdf, accessed 2026-10-17 (via `research/02-industries.md`) Source: Federal Maritime Commission, Detention and demurrage, https://www.fmc.gov/detention-and-demurrage/, accessed 2026-10-17 (via `research/02-industries.md`)

Cluster 3: Carrier invoice disputes and rate confirmations, priority 2 (shipper lens)

"Rate confirmation" (3.34x anchor) is the surprise of the pull. It is the first document our rules engine reads, and people search for it as a document in its own right (templates, required contents, whether it binds). Content that explains which rate-confirmation clauses decide a detention or accessorial dispute serves both sides of the market and links to the evidence-packet example. "Freight dispute" (0.10x) and "freight overcharge" (below threshold) are small. "Freight claim" (1.51x) looks large, but its related queries show it is about cargo loss and damage with FedEx and UPS, so we should not target it, and we should keep "claim" out of page titles where a searcher could mistake us for a cargo-claims service.

Keyword	Intent	Trends ratio	Monthly volume	Difficulty (qualitative)	Target page
rate confirmation / what is a rate confirmation	Informational	3.34	UNKNOWN — needs primary research	Medium: broker and load-board content likely	`/guides/rate-confirmation-detention-clauses/`
carrier invoice dispute / freight invoice dispute	Informational, shipper	freight dispute 0.10	UNKNOWN — needs primary research	Low-medium	`/guides/dispute-a-carrier-invoice/`
freight claim (cargo loss and damage)	Wrong intent: do not target	1.51	UNKNOWN — needs primary research	n/a	none

Cluster 4: Freight audit (vendor category), priority 2

This is the category investors recognise and searchers barely use in its long form. "Freight audit" carries modest interest, and its related queries ("freight audit and payment", "freight audit services") point at the outsourced audit-and-pay incumbents. Difficulty is high for the head term: software directories (Capterra, GetApp), listicles (Gitnux, Guideflow) and vendor blogs (FarEye, Hyperbots) fill the sampled results, and enterprise firms have years of domain authority. We should aim at the long tail where company size and job are specific ("freight audit for mid-market shippers", "contingency freight audit") with one neutral comparison page, published only after `research/01-competitors.md` has verified every competitor fact we would print.

Keyword	Intent	Trends ratio	Monthly volume	Difficulty (qualitative)	Target page
freight audit	Mixed, vendor-aware	1.00	UNKNOWN — needs primary research	High: directories, incumbents	home (secondary), `/freight-audit-alternatives/`
freight audit and payment / freight audit services	Vendor-aware (outsourced)	related query only	UNKNOWN — needs primary research	High	`/freight-audit-alternatives/`
freight audit software	Vendor-aware	0.02	UNKNOWN — needs primary research	High	`/freight-audit-alternatives/`
contingency freight audit / freight audit pricing	Decision	not pulled	UNKNOWN — needs primary research	Medium	`/pricing/`
freight invoice audit / freight bill audit	Vendor-aware	0.00 / 0.01	UNKNOWN — needs primary research	Medium	`/accessorial-audit/`

Source: Capterra, Freightview vs ShipScience comparison, https://www.capterra.com/compare/136895-176318/Freightview-vs-UPS-FedEx-Parcel-Audit, accessed 2026-10-17 Source: Gitnux, Top 10 best freight auditing software of 2026, https://gitnux.org/best/freight-auditing-software/, accessed 2026-10-17 Source: FarEye, Freight audit software (blog), https://fareye.com/resources/blogs/freight-audit-software, accessed 2026-10-17 Source: Guideflow, Best freight audit software tools, https://www.guideflow.com/blog/freight-audit-software, accessed 2026-10-17

Cluster 5: Detention billing (carrier side), priority 3: later and partner channel

Under the owner's decision this cluster waits: carriers become a later or partner channel. It keeps its pages in the plan because it has the most measurable demand, and because shippers receive carriers' detention invoices, so a shipper-side explainer of how detention is calculated serves both. "Detention pay" (2.33x anchor) and "detention charges" (1.24x) are the largest unpolluted freight terms we measured, and the evidence-packet idea answers the carrier's real question: how do I get paid for time I can prove? The difficulty is moderate rather than low because the people already writing for this query are well-funded adjacent businesses with their own motive to own it: carrier factoring companies (OTR Solutions), TMS vendors (Truckbase, Toro TMS) and driver-job boards (DriveMyWay). Their pages are explainers that end in a factoring or TMS pitch; none we sampled shows a worked calculation from a rate confirmation with the evidence attached, which is the gap a better page can fill.

Keyword	Intent	Trends ratio	Monthly volume	Difficulty (qualitative)	Target page
detention pay / detention pay trucking	Informational, carrier	2.33	UNKNOWN — needs primary research	Medium: factoring/TMS explainers rank	`/guides/detention-pay/`
detention charges / detention fee	Informational to commercial	1.24 / 0.36	UNKNOWN — needs primary research	Medium	`/guides/how-detention-is-calculated/`
how to bill detention / detention invoice	Commercial, carrier	not pulled	UNKNOWN — needs primary research	Medium	`/guides/bill-detention-with-proof/`

Keyword	Intent	Trends ratio	Monthly volume	Difficulty (qualitative)	Target page
detention billing software / detention tracking software	Vendor-aware	not pulled	UNKNOWN — needs primary research	Medium-high: TMS vendors and telematics marketplaces	`/detention-recovery/`
layover pay, TONU	Informational, glossary	0.71 / 2.77 (ambiguous)	UNKNOWN — needs primary research	Low-medium	glossary entries

Who publishes for this cluster today (web-search sample, 2026-10-17, not a Google results capture):

Source: OTR Solutions, Detention pay in trucking, https://otrsolutions.com/detention-pay-in-trucking/, accessed 2026-10-17 Source: Truckbase, Detention pay: what is it, https://www.truckbase.com/blog/dention-pay-what-is-it-how-to-get-your-truckers-money-worth, accessed 2026-10-17 Source: DriveMyWay, What is detention pay, https://www.drivemyway.com/blog/what-is-detention-pay-and-why-does-it-matter, accessed 2026-10-17 Source: Toro TMS, Driver detention management software: 7 top tools, https://www.torotms.com/blog/driver-detention-management-software, accessed 2026-10-17 Source: Geotab Marketplace, FleetOptos listing, https://marketplace.geotab.com/solutions/fleetoptos/, accessed 2026-10-17 Source: Platform Science Marketplace, DockClock, https://www.platformscience.com/marketplace/dockclock, accessed 2026-10-17

The software results matter for positioning. Trade-press coverage reports that McLeod and TMW (now Trimble) have offered detention modules inside their TMS products for years, and telematics-marketplace apps (FleetOptos on Geotab, DockClock on Platform Science, Detention Source on the App Store) automate the detention clock from geofences. A carrier searching "detention billing software" will meet those first. Our page has to say plainly that scoup.ai is neither a TMS nor a telematics app, and that its angle is the evidence packet built from the documents the carrier already holds.

`market-researcher` independently confirmed FleetOptos and Detention Source in `research/01-competitors.md` and added two more geofence-based tools that compete for this cluster: Hey Bubba (detention claims from ELD geofences inside an AI dispatcher; its `/pricing` page returned 404) and GeoStamp by WiseTech Global (geofence timestamps, listed on the Motive marketplace). The other publishers named above (DockClock, the McLeod and TMW modules, OTR Solutions, Truckbase, Toro TMS, DriveMyWay) come only from our search sample and are not yet in the competitor research. Neither of us has ranking or traffic data for any competitor.

Source: Hey Bubba, Detention management, https://bubba.ai/solutions/detention-management, accessed 2026-10-17 (via `research/01-competitors.md`)

Source: Bulk Transporter, McLeod module bills for downtime, https://www.bulktransporter.com/archive/article/21644416/wwwmcleodsoftwarecommodule-bills-for-downtime, accessed 2026-10-17 Source: Trucking Info, TMW adds detention tracking feature, https://www.truckinginfo.com/94877/tmw-adds-detention-tracking-feature, accessed 2026-10-17

Cluster 6: Entity and AI-answer terms (glossary)

Short definitional queries ("what is TONU", "what is a lumper fee", "what is detention in trucking") are the blocks AI answer engines and featured snippets tend to quote. Volumes are unknown and several terms are ambiguous in Trends, so we cover them as a glossary with one definition, one worked example and one primary source each, rather than as standalone pages. This carries over `seo-and-growth.md` Cluster D in substance.

4. What this means for the plan

The owner has chosen shippers and importers first, with carriers as a later or partner channel, and we order the plan that way. The search data does not argue for that choice and does not need to: segment choice rests on who pays, how many accounts exist and what the product already does (the accessorial check and the ocean

D&D module), which the research and finance files cover. Search has a narrower job, and the decision changes how we do it.

Four decisions follow, each reversible once real volume data arrives. First, shipper content leads, but it targets the plain words an AP team uses ("accessorial", "freight invoice", "freight bill") and the "how to dispute" question, not software phrases that barely register. Second, importer content leads alongside it and starts where importers actually search, with what a demurrage or detention charge means and how to check one, then moves to the FMC rule and the 30-day dispute window, where the ocean D&D module's checks give us something nobody else shows. Third, "freight audit" stays as a long-tail target for one neutral comparison page; the public category label is "detention, accessorial and D&D recovery for the shipper side". Fourth, carrier detention pages move to the back of the queue. They are the largest measured demand we found, so we keep them in the plan for the partner channel and for shippers who receive detention invoices, but no carrier page displaces a shipper or importer page.

We should be honest about the cost. Leading with the smaller-search segments means organic discovery will be slower than a carrier-first plan would have been, if search were the channel that mattered. It is not, at this stage: design partners come from outreach, and the shipper and importer pages exist to convert those conversations.

Two traps are worth writing down so nobody walks into them later. "Freight claim" looks like a large opportunity and is the wrong job (cargo loss and damage). "Detention time" and "TONU" look large and are polluted by other meanings, so any future volume figure for them must be filtered to freight intent before it informs a decision.

The investor-facing reading is equally plain. Search will not be a primary acquisition channel at this stage; the chosen buyers rarely search for software at all. SEO earns its place by making the company credible when a design partner looks us up, and by owning a handful of worked-example pages that nobody else has built.

5. How to replace UNKNOWN with real numbers

We would close the volume gap in this order, cheapest first. Google Search Console gives real impressions and queries for our own pages from the day a verified property exists; it measures our visibility, not market volume, but it is free and exact. Google Keyword Planner gives market volume, but trade press reports that low-spend accounts see banded ranges instead of exact figures; a small, time-boxed search campaign on the Cluster 1 and 2 phrases would unlock exact figures and test messaging at the same time. A one-month paid SEO tool subscription (Semrush, Ahrefs or Moz) would add difficulty scores and competitor-ranking exports; we have not priced it and record no figure here. The qualitative difficulty calls above should be replaced the moment any of these exist.

Source: Search Engine Land, Google officially throttling Keyword Planner data for low spending AdWords accounts, https://searchengineland.com/google-officially-throttling-keyword-planner-data-low-spending-adwords-accounts-255795, accessed 2026-10-17 (reports banded ranges such as 100-1K for low-spend accounts; Google never defined the spend threshold)

6. Open questions

These items are `UNKNOWN — needs primary research` and block a quantitative keyword plan:

- Absolute US monthly search volume for every keyword in section 3 (no tool account; no Keyword Planner spend).
- Keyword difficulty scores from any tool; current calls are qualitative.
- Freight-intent share of "detention time", "TONU" and "demurrage" (Trends cannot separate meanings).
- Who actually ranks on a captured Google results page (our competitor sample comes from a web-search API).

- Competitors' own organic traffic and ranking keywords. `market-researcher` named the domains to export once a tool exists: loop.com, cassinfo.com, traxtech.com, ctsi-global.com, corporate.nvisionglobal.com, intelligentaudit.com, bluecargo.io, bubba.ai (Hey Bubba, carrier detention capture; competes directly for Cluster 1), the FleetOptos Geotab listing, and portcast.io (adjacent).
- Federal Register document 2025-23920 (removal of 46 CFR 541.4) read first-hand, not via the build's source comment (bot-check and unreadable PDF on 2026-10-17).
- Search interest in the category label marketing uses publicly, "detention and accessorial recovery" (and "detention recovery"): not in the 2026-10-17 pull; add it to `data/trends_pull.py` on the next run. We expect it to sit below the Trends threshold, which is fine for a label used on our own pages.
- Norwegian or EU keyword demand: out of scope for this step; the buyers in the business plan are US.

Content plan: pillars, clusters and a 12-month calendar

> **Status.** scoup.ai is pre-product: synthetic data only, no customers, no live product, never sold. Nothing in this plan has been published or measured. The calendar is a capacity plan for one founder plus a freight-domain reviewer we have not yet recruited; it is not a traffic forecast and carries no traffic, ranking or lead targets. Keyword evidence and its limits are in `keyword-research.md`. This plan builds on `marketing/seo-and-growth.md` sections 3 to 5 and on the site map in `marketing/landing/outline.md`. Re-ordered on 2026-10-17 for the owner's shipper- and importer-first decision; carriers are a later or partner channel.

1. What content has to do for us right now

The owner decided on 2026-10-17 that scoup.ai goes to market shipper- and importer-first, with carriers as a later or partner channel. This plan is ordered by that decision. Our go-to-market still starts with design-partner outreach, not search, because nobody has yet confirmed on real closed files that the leakage we model is real. Content therefore has three jobs, in this order. It must make the company credible when a shipper's AP lead or an importer's logistics manager looks us up after an outreach email. It must give the founder specific things to send in that outreach: an annotated sample, a checklist, a line-by-line dispute walkthrough. And it must own a few pages nobody else has built, such as an invoice line checked against its contract clause with the arithmetic shown, or an ocean D&D invoice checked element by element against the FMC rule. Volume of posts is not a job.

The keyword pull tells us how to write for these buyers rather than whether to. Software phrases a shipper might type ("freight audit software", "freight invoice audit") sit at or below the Google Trends threshold, but the plain words an AP team uses do register: "freight bill" at 3.88 times the interest of "freight audit", "freight invoice" at 1.47 times, "accessorial" at 0.88 times. Importers search for what a charge means ("demurrage" 4.88 times, broad; "demurrage charges" 0.22) and for carriers' tariffs far more than for the rule. Carrier detention terms remain the largest clean demand we measured ("detention pay" 2.33 times), which is why carrier pages stay in the plan at the back, ready for the partner channel.

Source: Google Trends, US, 2021-10-01 to 2026-09-30, `marketing/seo/data/trends_summary.csv` and `marketing/seo/data/trends_summary_F_importer_G_shipper.csv`, pulled 2026-10-17 (ratios are our arithmetic; see `keyword-research.md` section 2)

2. Pillars and clusters

Each pillar is a hub page with spokes, numbered in the owner's priority order. The content pillars map onto the three messaging pillars `marketing` settled in `marketing/brand/README.md` section 4: (1) proof attached, (2) a person decides, (3) built for the loads big audit programs skip. The public category label is "detention, accessorial and D&D recovery for the shipper side", not "freight audit"; "freight audit" stays only where a searcher uses it (Cluster 4). The shipper hero is marketing's workload line, "Take accessorial review off AP's desk. Every line checked against the rate confirmation, with the source attached" (`marketing/marketing-plan.md` section 2.3). Shipper page titles and descriptions lead with that workload framing where it fits the query (solution page, "how to check a freight invoice", the AP walkthrough). Dispute-intent pages ("how to dispute accessorial charges") keep titles that answer the dispute question directly, because that is what the searcher typed. "Stop paying accessorials that were never authorized" stays as the A/B alternate and as an H1 option for `/accessorial-audit/`. The reason for the switch is a broker's survey: RXO, a freight broker with no audit product but not neutral, reports that "61% of shippers agree that managing accessorial requests is one of the biggest time-wasters for their accounting team" (750 shippers and carriers, published 25 November 2024). Any page quoting it says it is a broker survey and carries this source line.

Source: RXO, What shippers and carriers really think about trucking accessorials (original research, 25 November 2024), https://rxo.com/resources/research/freight-accessorials/, accessed 2026-10-17 (via `research/04-gaps.md` gap 5)

The importer hero line is "Check every D&D invoice before the 30-day window closes", supported by "We check each container detention and demurrage invoice against the FMC billing rule's content and timing requirements, and list what to dispute and by when" (`marketing/marketing-plan.md` section 2.3). Every page that describes the rule carries the Thompson Hine source and the note "confirm with counsel or your customs broker". Shipper pages speak to the transportation or logistics manager (champion), the AP manager or controller (user) and the CFO or VP finance (buyer); importer pages to the import or logistics manager and AP. Carriers are a later channel reached through factoring and broker partners, so carrier pages (P5) are written to be forwarded by those partners. The keyword clusters are the ones in `keyword-research.md` section 3, renumbered there on the same date (Cluster 1 shipper, Cluster 2 importer, Cluster 5 carrier).

"Proof attached" leads every page. "A person decides" is the answer to the trust objection ("will you dispute lines on our behalf?"), placed in the body and FAQ, not in titles or H1s, because `research/04-gaps.md` found no outside demand for approval controls. Wherever the ATRI figures appear, the page says they measure the cost of waiting, not recoverable dollars. Importer pages may describe the ocean D&D module only as pre-product software tested on synthetic invoices, with no accuracy or recovery claim.

Pillar (hub URL)	Promise to the reader	Keyword clusters	Messaging pillar	Primary audience	Role
P1 Accessorial review, off AP's desk (`/guides/accessorials/`)	How to check each freight invoice line against the rate confirmation, and dispute the ones that do not match	Cluster 1, Cluster 6	Proof attached (trust answer: a person decides)	Mid-market shipper AP and logistics	Lead segment; design-partner conversion
P2 Demurrage and detention, checked (`/guides/demurrage-detention/`)	What a D&D charge means, what the FMC rule requires on the invoice, and how to dispute inside 30 days	Cluster 2	Proof attached	Small and medium importers, their forwarders	Lead segment; authority and conversion
P3 The contract is the evidence (`/guides/rate-confirmation/`)	Which rate-confirmation and tariff clauses decide an accessorial or D&D dispute	Cluster 3	Proof attached	Shippers and importers	Bridge between P1 and P2
P4 Choosing a recovery approach (`/freight-audit-alternatives/`, `/pricing/`)	Honest comparison of in-house, TMS module, outsourced audit (Loop and Intelligent Audit's Catalyst contest small shippers), ocean D&D tools (BlueCargo is strong; we concede it) and us, conceding incumbents' strengths first; our angle is the per-line evidence packet, contingency pricing and one workflow for truck detention and container D&D on the same shipments (dated sources in `marketing/marketing-plan.md` section 2.2)	Cluster 4	Built for the loads big audit programs skip	Shippers and importers evaluating vendors	Decision support

Pillar (hub URL)	Promise to the reader	Keyword clusters	Messaging pillar	Primary audience	Role
P5 Detention you can prove (`/guides/detention/`)	How truck detention is calculated, billed and paid, with the arithmetic	Cluster 5, Cluster 6	Proof attached	Carriers (later and partner channel); shippers receiving detention invoices	Deferred; partner channel

Every spoke follows the explainer template in `marketing/landing/outline.md`: the question as H1, a direct answer in the first 100 words, a worked example with arithmetic, "what varies by contract", primary sources with dates, a named author and reviewer (reviewer named only with consent), and one next step. Worked examples use synthetic loads and synthetic ocean invoices from the repository's fixtures and say so.

Internal linking is planned, not left to chance. Every spoke links up to its hub and across to the matching solution page: `/accessorial-audit/` for P1 and `/demurrage-detention-audit/` for P2 (a new URL; the landing outline had no importer page). P3 links to both. The sample evidence packet is linked from every P1 and P2 page, because it is what a design-partner conversation is built around. Carrier pages (P5) link back into P1, since a shipper is the party who receives a carrier's detention invoice.

3. Twelve-month calendar (November 2026 to October 2027)

The calendar assumes about two substantial pieces a month, which is what one founder can write and get reviewed while running outreach. Nothing publishes until the domain is secured and the site fixes in `technical-checklist.md` are done; that hold stands. Quarter one builds the shipper and importer foundations, quarter two the dispute playbooks and checklists outreach needs, quarter three the decision pages and original-data pieces, and carrier pages wait until the partner channel opens. A "DP" mark means the piece is built to be sent in design-partner outreach or to convert a visitor into a design-partner conversation. Any month can slip; the order matters more than the dates.

Month	Piece	URL	Pillar / cluster	Format	DP	Gate before publishing
Nov 2026	Site fixes, domain decision, Search Console property	n/a	n/a	Technical		`technical-checklist.md` P0 items done; domain secured
Nov 2026	Sample evidence packets: one shipper accessorial invoice, one ocean D&D invoice (synthetic)	`/evidence-packet-example/`	P1/P2	Annotated samples	DP	Every figure traced to a fixture; "synthetic" label on every page
Dec 2026	Shipper solution page	`/accessorial-audit/`	P1 / C1	Solution page	DP	No capability claim beyond the prototype
Dec 2026	How to dispute accessorial charges, line by line	`/guides/dispute-accessorial-charges/`	P1 / C1	Playbook + checklist	DP	Domain-reviewer pass
Jan 2027	Importer solution page	`/demurrage-detention-audit/`	P2 / C2	Solution page	DP	Module described as pre-product, synthetic data only

Month	Piece	URL	Pillar / cluster	Format	DP	Gate before publishing
Jan 2027	Demurrage vs detention: what each charge means and who bills it	`/guides/demurrage-vs-detention/`	P2 / C2	Explainer		Reviewer pass
Feb 2027	FMC D&D invoice checklist (printable)	`/guides/fmc-dd-invoice-checklist/`	P2 / C2	Checklist (section 5, item 3)	DP	Current 46 CFR 541 text and FR doc 2025-23920 read first-hand; counsel review
Feb 2027	How to check a freight invoice before you pay it	`/guides/how-to-check-a-freight-invoice/`	P1 / C1	AP walkthrough	DP	Reviewer pass
Mar 2027	How to dispute demurrage charges inside the 30-day window	`/guides/dispute-demurrage-charges/`	P2 / C2	Playbook	DP	Same rule check as February
Mar 2027	Glossary hub (accessorial, lumper, demurrage, detention, per diem, free time, TONU)	`/guides/glossary/`	C6	Glossary		One primary source per entry
Apr 2027	Rate-confirmation and tariff clauses that decide accessorial and D&D disputes	`/guides/rate-confirmation-detention-clauses/`	P3 / C3	Annotated clause guide	DP	Reviewer pass; no third-party templates copied
Apr 2027	Fuel surcharge check: EIA diesel index and your FSC table	`/guides/fuel-surcharge-check/`	P1	Original-data tool (section 5, item 2)	DP	EIA reuse policy read; weekly update job tested
May 2027	Lumper fees: when they are legitimate, when to flag them	`/guides/lumper-fees/`	P1 / C1	Explainer		Reviewer pass
May 2027	How to dispute a detention charge (shipper side)	`/guides/dispute-a-detention-charge/`	P1 / C1	Playbook	DP	Reviewer pass
Jun 2027	Ocean D&D free-time tracker (public carrier tariffs)	`/guides/dd-free-time-tracker/`	P2 / C2	Original-data tracker (section 5, item 6)		Feasibility and terms check per carrier site
Jun 2027	Recovery calculator, shipper and importer variants (spec: `marketing/landing/calculator-spec.md` plus marketing's draft variants in `marketing/brand/calculator-variants.md`, pending merge)	`/calculator/`	P1/P2	Tool	DP	Unit test of the worked example; CSP change reviewed (it needs script)
Jul 2027	Freight audit and D&D audit alternatives for mid-market shippers and importers	`/freight-audit-alternatives/`	P4 / C4	Neutral comparison	DP	Every competitor fact verified from their own pages that week, date-stamped
Jul 2027	Pricing: the contingency model explained	`/pricing/`	P4 / C4	Decision page	DP	Only after a real offer exists; no fee number before that
Aug 2027	What detention and D&D cost, and why cost is not recoverable dollars	`/guides/detention-cost-data/`	P1/P2	Original-data explainer (section 5, item 1)		ATRI, DOT OIG and FMC figures re-checked against live sources

Month	Piece	URL	Pillar / cluster	Format	DP	Gate before publishing
Aug 2027	How truck detention is calculated: clock start, free time, rounding, cap (shipper lens)	`/guides/how-detention-is-calculated/`	P5 / C5	Worked example		Reviewer pass
Sep 2027	Building in public: what design partners taught us (no metrics unless consented)	`/notes/…`	P1/P2	Founder note	DP	Partner consent for any detail
Sep 2027	Refresh pass on every page older than six months	n/a	all	Maintenance		Search Console data, if any
Oct 2027	Carrier solution page, only if the partner channel has opened	`/detention-recovery/`	P5 / C5	Solution page		Owner go-ahead on the carrier channel
Oct 2027	Re-plan from twelve months of real Search Console and outreach data	n/a	n/a	Review		Keyword volumes replaced where measured

Moved out of the twelve months by the owner's decision, and kept ready for the carrier partner channel: detention pay for carriers (`/guides/detention-pay/`), detention and the 14-hour clock (section 5, item 4), and detention rule variants computed side by side (section 5, item 5). Nothing in the calendar depends on customer data. If a design partner later consents to an anonymised case study, it slots in as an extra piece; it never replaces a gate.

4. Which pieces win design partners, and how they are used

Design partners will come mostly from outreach and warm introductions, so the DP pieces are built to be attached to a message rather than found. Each answers one objection a shipper or importer raises before agreeing to share closed invoices. The sequence below is how we expect the founder to use them; it is a hypothesis to test in the first ten conversations, and the order will change with what partners actually ask.

Objection a prospect raises	Piece that answers it	How it is used
"What would I actually get?"	Sample evidence packets (accessorial and ocean D&D, synthetic)	Linked in the first outreach email; the most concrete thing we have
"Which invoice lines would you even question?"	How to check a freight invoice; dispute accessorial charges; contract clauses guide	For shipper AP leads; shows the line-against-clause check
"Our D&D invoices come from big carriers; what is there to check?"	FMC D&D invoice checklist; dispute demurrage charges	For importers; the checklist mirrors the module's field-by-field checks
"Is there money in this for us?"	Recovery calculator (shipper and importer variants)	Prospect enters their own numbers; output is a range with the arithmetic, never a promise
"We already have an audit firm / a D&D tool"	Alternatives page; shipper and importer solution pages	Honest about where incumbents are stronger; positions the evidence packet as the gap
"What happens to our documents?"	Existing `/security/` page plus the draft handling terms	Before any file changes hands; nothing real is accepted until the controls on that page exist

The pieces without a DP mark (glossary, D&D explainer, free-time tracker, cost explainer, carrier pages) exist for search authority, links and the later carrier channel. They should not displace a DP piece when time runs short.

5. Original-data content from public sources (no customer data)

These ideas produce something nobody else publishes in this exact form, using only public data and our own rules engine run on synthetic inputs. Each lists its source, what we would compute, and the honesty limit it must carry. None requires a design partner's data, which is why they can start before the first pilot. Under the

shipper- and importer-first decision the order of use is items 3, 6, 2 and 1 in the twelve months; items 4 and 5 are carrier pieces held for the partner channel; item 7 can go any time.

1. What detention costs, separated from what is recoverable. ATRI's September 2024 research reports that "the trucking industry lost $3.6 billion in direct expenses and $11.5 billion in lost productivity from driver detention in 2023", that drivers were detained at 39.3 percent of stops, and that while 94.5 percent of fleets charge detention they are paid on fewer than 50 percent of invoices. The two dollar figures sum to $15.1 billion (our arithmetic: 3.6 + 11.5). The US DOT Office of Inspector General's January 2018 report estimated a $1.1 to $1.3 billion annual earnings loss for for-hire truckload drivers and said accurate industry-wide detention data do not exist. The piece explains, with arithmetic, why lost productivity is not a pool of recoverable invoices, which is the distinction `seo-and-growth.md` already insists on.

Source: American Transportation Research Institute, New research documents substantial financial and safety impacts from truck driver detention (10 September 2024), https://truckingresearch.org/2024/09/new-research-documents-substantial-financial-and-safety-impacts-from-truck-driver-detention/, accessed 2026-10-17
Source: US DOT Office of Inspector General, Estimates show commercial driver detention increases crash risks and costs, but current data limit further analysis (ST2018019, 31 January 2018), https://www.oversight.gov/reports/audit/estimates-show-commercial-driver-detention-increases-crash-risks-and-costs-current, accessed 2026-10-17

2. Fuel surcharge check against the EIA diesel index. Where a contract's fuel surcharge table references the EIA weekly US on-highway diesel price (how common that is: `UNKNOWN — needs primary research`), the index is public and EIA releases it on Tuesdays (latest shown on 2026-10-17: $6.179 per gallon for the week of 12 October 2026, released 14 October). A page that recomputes the surcharge a table implies for any week, from the published index, lets a shipper check an FSC invoice line in a minute. EIA data are free through its open-data API with a free key and are subject to EIA's copyright and reuse policy, which we must read before shipping. The TS build (branch `feat/hardening`) already has a fuel-surcharge check (`api/src/intelligence/fuel.ts`) that compares an invoice's FSC line with a tenant's surcharge table and an imported diesel price series; it runs on synthetic data only, so the page may present it as a pre-product feature with no accuracy claim. (An earlier draft said no such rule existed; that was wrong.)

Source: US Energy Information Administration, Gasoline and diesel fuel update, https://www.eia.gov/petroleum/gasdiesel/, accessed 2026-10-17 Source: US Energy Information Administration, Open data (API), https://www.eia.gov/opendata/, accessed 2026-10-17

3. FMC demurrage and detention invoice checklist. The FMC's 2024 rule (46 CFR part 541) sets minimum information for D&D invoices, a 30-day invoicing deadline for ocean carriers and terminal operators, and at least 30 days for the billed party to request mitigation, refund or waiver. A printable checklist mapping each required element to "present or missing" on an invoice is useful and checkable. It mirrors the field-by-field checks in the build's ocean D&D module (`packages/shared/src/ocean-dd/`), so it doubles as an honest preview of that feature, labelled pre-product and synthetic. Since a September 2025 D.C. Circuit ruling vacated only 46 CFR 541.4, the content and dispute provisions the checklist uses remain in force; Federal Register document 2025-23920 must still be read first-hand.

Source: HFW, FMC's new demurrage and detention regulations become effective May 2024, https://www.hfw.com/insights/fmcs-new-demurrage-and-detention-regulations-become-effective-may-2024, accessed 2026-10-17

4. Detention and the 14-hour clock. The federal hours-of-service rule for property carriers (49 CFR part 395) limits driving to a 14-hour on-duty window (our reading of the rule; the FMCSA summary page returned HTTP 403 on 2026-10-17, so the text must be checked on eCFR before publishing). A worked example on a synthetic load shows how many driving hours a three-hour detention removes and what that means for the next load. We cite the regulation text directly and do not model safety outcomes ourselves; the DOT OIG crash estimate above is quoted as theirs.

5. Detention rule variants computed side by side. Our rules engine already implements clock start (arrival vs appointment), free time, rounding and caps. Running one synthetic arrival and departure through a dozen common contract patterns and printing every result with its arithmetic produces a reference table nobody publishes. It is our own computation on synthetic data and is labelled that way; a domain reviewer confirms the patterns are realistic.

6. Ocean D&D free-time tracker. Google Trends shows "msc detention and demurrage tariff" as a rising related query (+650 percent over the window), so importers look for carrier tariffs. A tracker that records each major ocean carrier's published free days per port, with the source page and the date we read it, would be original and linkable. It is the most effort and the furthest from the product, so it sits last, after a check of each carrier site's terms of use.

Source: Google Trends related queries for "demurrage", US, 2021-10-01 to 2026-09-30, `marketing/seo/data/trends_related.csv`, pulled 2026-10-17

7. What shippers, importers and carriers search for. A short public note on our own Trends pull (method, script, the "freight claim" trap, the gap between carrier vocabulary and software vocabulary) is honest, cheap and useful to other people in the niche. It reports relative interest only.

6. Rules every piece follows

These rules come from the project's honesty constraints and from what Google has said about structured data. Every number is either cited to a primary source with a date or labelled as our arithmetic or an illustration. No page claims a recovery rate, an accuracy figure, a customer, a certification or an integration, because none exists. A freight-domain reviewer checks every P1, P2 and P3 piece before it goes live; until we have recruited one, those pieces stay in draft. Comparison content uses only facts competitors publish about themselves, re-checked the week we publish. We mark up FAQ sections honestly but expect no FAQ rich result: since August 2023 Google shows FAQ rich results only for well-known government and health sites.

Source: Google Search Central Blog, Changes to HowTo and FAQ rich results (8 August 2023), https://developers.google.com/search/blog/2023/08/howto-faq-changes, accessed 2026-10-17

7. How we will know whether it works

We set no traffic targets, because there is no baseline and no measured volume. The measures that matter in year one, in order, are: design-partner conversations where the prospect had opened a DP piece (tracked by asking, and by tagged links in outreach); Search Console impressions and queries per pillar once the site is verified on its final domain; and which pieces prospects forward inside their own company. Numeric targets get written down only after two months of real data, and they are written as targets, never as results.

8. Open questions

- Domain: scoup.ai did not resolve on 2026-10-17 and the .ai registry's RDAP service returned "Object not found"; publishing content before the final domain is secured means rebuilding search equity later. See `technical-checklist.md` item T1.
- A freight-domain reviewer: none recruited; every P1 to P3 piece is blocked on one.
- Merge of marketing's shipper and importer calculator variants (`marketing/brand/calculator-variants.md`) into `marketing/landing/calculator-spec.md`; the owner of `marketing/landing/` has not been assigned.
- Share of fuel surcharge tables keyed to the EIA index: `UNKNOWN — needs primary research`.
- First-hand reading of Federal Register document 2025-23920 and of 49 CFR part 395 (blocked or unreadable on 2026-10-17).
- Absolute search volumes for every target keyword: `UNKNOWN — needs primary research` (see `keyword-research.md`).

Technical SEO checklist: audit of `marketing/site/`

> **Status.** Desk audit of the static marketing site source plus a few read-only requests to the copy that is already public. No edits were made to the site. scoup.ai is pre-product: no customers, no traffic data, no Search Console property we know of, so nothing here measures rankings or field performance. Every fix below names the file to change; the site's owner (not `seo`) makes the changes.

1. What was audited, and how

We audited three things, because they are not the same site. The **master source** in `marketing/site/` (commit 5828944 on master, copied into this worktree) still brands the site "Freight Recovery". The **brand branch** `feat/scoup-brand` (commits 98c9a54 and 3933aae) rebrands to scoup.ai, adds an Open Graph card, a favicon and base-path support, and adds `.github/workflows/pages.yml`, which deploys the site to GitHub Pages. The **live site** answers at https://socrates6969.github.io/freight-recovery/ with HTTP 200 and the scoup.ai title, and its sitemap shows `lastmod` 2026-10-12, so a build from the brand work is public today even though the master `marketing/site/README.md` still says "Not deployed".

We built both source versions in a scratch directory with Node 24.21.0 and ran the site's own `npm test`; both passed 9 of 9 page checks (master with `SITE_URL=https://www.scoup.ai`, brand with the github.io URL). We read the generated HTML, `sitemap.xml` and `robots.txt`, fetched the live response headers, and compared the result against Google Search Central documentation. File paths below are relative to the repository root and refer to master unless marked "brand".

The overall picture is good for a site this young. The builder already does most things right: one H1 per page, unique titles and descriptions with length warnings, canonical tags, Open Graph and Twitter tags, honest JSON-LD with no ratings and no `SoftwareApplication`, a strict CSP, no JavaScript and no third-party requests, and a test suite that fails the build on banned claims. The problems are mostly about where the site is hosted, what is indexable, and what is missing for search to have anything to rank.

2. Findings summary

ID	Priority	Finding	Main file(s)
T1	P0	Site is public on a github.io sub-path; final domain appears unregistered	`.github/workflows/pages.yml` (brand), `marketing/site/build.mjs`
T2	P0	Security headers in `_headers` are not served on GitHub Pages	`marketing/site/static/_headers`, `marketing/site/README.md`
T3	P1	Draft legal pages are indexable and listed in the sitemap	`marketing/site/pages/privacy.html`, `terms.html`, `cookies.html`
T4	P1	Every sitemap `lastmod` is the build date	`marketing/site/build.mjs`
T5	P1	No social image or favicon on master	`marketing/site/build.mjs` (fixed on brand)
T6	P1	Entity name "Freight Recovery (pre-product)" in schema and `og:site_name`	`marketing/site/build.mjs` (fixed on brand)
T7	P1	No page targets a measured keyword cluster; generic H1 on how-it-works	`marketing/site/pages/*.html`
T8	P2	Organization schema has no logo; WebPage schema has no dates	`marketing/site/build.mjs`
T9	P2	FAQPage markup will not produce a rich result	`marketing/site/pages/index.html`
T10	P2	404 page carries a canonical tag	`marketing/site/build.mjs`

ID	Priority	Finding	Main file(s)
T11	P2	Obsolete `interest-cohort` directive; HSTS without `includeSubDomains`	`marketing/site/static/_headers`
T12	P2	Personal mailbox published as the contact address	`.github/workflows/pages.yml` (brand)
T13	P2	No skip link; minor accessibility gaps	`marketing/site/build.mjs`, `static/style.css`
T14	P2	Tests do not guard the new rules (noindex drafts, lastmod, og:image)	`marketing/site/test.mjs`
T15	P3	Performance: nothing to fix now; measure field data after launch	n/a

3. Findings and fixes

T1 (P0). Public on a github.io sub-path, and the final domain looks unregistered

The brand workflow builds with `SITE_URL: https://socrates6969.github.io/freight-recovery`, so every canonical URL, `og:url`, JSON-LD `@id` and sitemap entry points at a GitHub-owned hostname under a repository path. Three problems follow. Search equity earned now accrues to a URL we will abandon. Google ignores a `robots.txt` that is not at the host root ("Crawlers don't check for robots.txt files in subdirectories"), and the root https://socrates6969.github.io/robots.txt returned HTTP 404 on 2026-10-17, so the site's own `robots.txt` and its `Sitemap:` line are invisible to crawlers. And Search Console can only verify a URL-prefix property for a sub-path, not a Domain property.

The brand guidelines name the company and product "scoup.ai", but on 2026-10-17 `scoup.ai` did not resolve in DNS and the .ai registry's RDAP service (`rdap.identitydigital.services`) answered `404 Object not found` for it. That strongly suggests the domain is not registered. We did not check trademark or availability beyond that one lookup.

Source: Google Search Central, Create and submit a robots.txt file, https://developers.google.com/search/docs/crawling-indexing/robots/robots_txt, accessed 2026-10-17

Fix, in order: (1) the owner decides and secures the production domain (scoup.ai or another) before more public content ships; (2) point the host at it and rebuild with `SITE_URL=https://<domain>` in `.github/workflows/pages.yml` (brand) or the new host's settings, so the builder's existing logic rewrites canonicals and the sitemap; (3) confirm whether the old github.io URLs 301 to the new domain (GitHub's custom-domain documentation we read does not state it for project sites); if they do not, keep them out of search by taking that deployment down once the new one is live; (4) verify a Domain property in Search Console by DNS and submit `https://<domain>/sitemap.xml`. Until step 1 is done, consider adding `"noindex": true` site-wide via a build flag so the github.io copy is not indexed at all.

T2 (P0). The security headers are not served

`marketing/site/static/_headers` holds the CSP with `frame-ancestors 'none'`, `X-Content-Type-Options`, `X-Frame-Options`, `Referrer-Policy`, `Permissions-Policy`, COOP and HSTS. That file format is honoured by Netlify and Cloudflare Pages, not by GitHub Pages. The live response on 2026-10-17 carried only `Server: GitHub.com`, `Strict-Transport-Security: max-age=31556952` and `Cache-Control: max-age=600`: no CSP header, no `X-Frame-Options`, no `nosniff`. The meta CSP in `build.mjs` still applies to scripts, styles and images, but a meta CSP cannot set `frame-ancestors`, so the pages can be framed by any site today. This is mainly a security item; it matters to SEO because the site's `/security/` page and README promise these headers, and a reviewer who checks will find the gap.

Source: Cloudflare, Pages headers configuration, https://developers.cloudflare.com/pages/configuration/headers/, accessed 2026-10-17 Source: GitHub Community, HSTS (and other security headers) for GitHub Pages with custom domains (discussion 4444), https://github.com/orgs/community/discussions/4444, accessed 2026-10-17

Fix: host on a platform that applies `_headers` (Cloudflare Pages or Netlify; no code change needed) or put a CDN in front of GitHub Pages that sets them; then confirm with `curl -sI`. Update the deploy section of `marketing/site/README.md` so it no longer implies the headers apply on every host.

T3 (P1). Draft legal pages are indexable

`pages/privacy.html`, `pages/terms.html` and `pages/cookies.html` say "DRAFT - must be reviewed by a licensed Norwegian advokat before use" and "not in force", yet their meta has no `noindex`, so `build.mjs` emits `<meta name="robots" content="index, follow">` and adds them to `sitemap.xml`. The live `/privacy/` page shows exactly that. A search result for "scoup.ai privacy policy" landing on a document that says it is not in force helps nobody. Fix: add `"noindex": true` to the meta line of each of the three pages; the builder already drops `noindex` pages from the sitemap. Remove the flag when the advokat has signed off.

T4 (P1). Sitemap `lastmod` is the build date for every URL

`build.mjs` writes `<lastmod>${today}</lastmod>` for every page, so each rebuild claims every page changed. Google says it uses `lastmod` only "if it's consistently and verifiably ... accurate", so this teaches Google to ignore it. Fix: add an `"updated": "YYYY-MM-DD"` field to each page's meta line and write that value (fall back to the file's last git commit date via `git log -1 --format=%cs -- <file>` if absent); reuse the same date for `dateModified` in T8. Google also ignores `priority` and `changefreq`, so do not add them.

Source: Google Search Central, Build and submit a sitemap, https://developers.google.com/search/docs/crawling-indexing/sitemaps/build-sitemap, accessed 2026-10-17

T5 (P1). No social image or favicon on master

The master `layout()` in `build.mjs` emits `twitter:card` "summary" and no `og:image`, and there is no favicon, so a shared link renders as bare text and browser tabs and some result displays show a generic icon. The brand branch fixes all three: `static/og-card.png` (57,555 bytes, 1200 × 630, rendered from `og/og-card.html`), `og:image` with width, height and alt, `summary_large_image`, and `favicon.svg` copied from `brand/`. Fix: merge the brand branch's `build.mjs`, `static/og-card.png` and `og/` into master rather than re-implementing them.

T6 (P1). The entity name in structured data is wrong on master

Master sets `SITE_NAME = "Freight Recovery (pre-product)"`, which becomes the `Organization` and `WebSite` name in JSON-LD and `og:site_name`. A status note in an entity name is the kind of string Google may display as the site name. The brand rules say the name is always lowercase "scoup.ai", and the brand branch sets `SITE_NAME = "scoup.ai"` and appends " – scoup.ai" to titles (all eight brand titles stay at or under 59 characters; we measured them). Fix: take the brand branch value; keep the pre-product status in the visible banner and the Organization `description`, where it already is.

T7 (P1). Nothing on the site targets a measured keyword cluster

The site has eight indexable pages: home, how it works, security, about, contact and three draft legal pages. None targets the carrier detention, accessorial dispute or rate-confirmation clusters in `keyword-research.md`, and the pages that would (`/detention-recovery/`, `/accessorial-audit/`, `/guides/...`, `/evidence-packet-example/`) are listed as deferred in `marketing/site/README.md`. The how-it-works page has a good title ("How freight detention recovery works") but a generic H1 ("How it works"). Fix: change the H1 in `pages/how-it-works.html` to match the title's topic; then add pages in the order of `content-plan.md` section 3, each as a fragment in `pages/` with its own meta line. Add the new pages to `NAV` in `build.mjs` only when they exist; keep the nav short.

T8 (P2). Structured data is thin where it is cheap to improve

`jsonLd()` in `build.mjs` emits Organization, WebSite, WebPage, BreadcrumbList and (home only) FAQPage, all honest. Two gaps: Organization has no `logo`, and WebPage has no `datePublished` or `dateModified`. Fix: add `logo` pointing to a brand image that meets Google's Organization structured-data image requirements (check format and minimum size in that documentation before choosing the file), and add `dateModified` from the T4 `updated` field. Add `sameAs` only for profiles that really exist. When guides ship, give them `Article` with a named author; never add `Review`, `AggregateRating` or `SoftwareApplication` until something real is available.

T9 (P2). FAQPage markup is honest but will not earn a rich result

The home page FAQ is visible in `<details>` elements and mirrored in `FAQPage` JSON-LD, which the test suite enforces. Since August 2023 Google shows FAQ rich results only for well-known, authoritative government and health sites, and says there is no need to remove existing markup. Fix: none required; just do not count on the rich result, and do not add FAQ blocks to pages for SEO reasons.

Source: Google Search Central Blog, Changes to HowTo and FAQ rich results (8 August 2023), https://developers.google.com/search/blog/2023/08/howto-faq-changes, accessed 2026-10-17

T10 (P2). The 404 page carries a canonical tag

`layout()` adds `<link rel="canonical">` and `og:url` to every page, including `pages/404.html`, whose path is `/404.html` and which is correctly `noindex`. A canonical on an error page is meaningless at best. Fix: in `build.mjs`, skip the canonical and `og:url` tags when `meta.noindex` is true and the path is `/404.html`, and relax the matching canonical check in `test.mjs` for that page.

T11 (P2). Header details

In `static/_headers`, `Permissions-Policy` lists `interest-cohort=()`, which belongs to Google's withdrawn FLoC proposal and does nothing useful now; remove it to keep the policy clean. Once the site is on its own domain, add `includeSubDomains` to HSTS if every subdomain will serve HTTPS. `X-Frame-Options: DENY` duplicates `frame-ancestors 'none'` but is harmless for older browsers. When the calculator ships it needs a script, so the CSP in both `build.mjs` and `_headers` will need `script-src 'self'` for that page; `test.mjs` currently fails any non-JSON-LD script and will need a scoped exception.

T12 (P2). A personal mailbox is the public contact address

The brand workflow sets `CONTACT_EMAIL: mariuscarlsson@gmail.com`, which `build.mjs` writes into the footer and the contact page of the public site. For a company asking businesses to share freight documents, a personal free-mail address undercuts the trust the security page works to build, and it exposes the founder's inbox to scraping. Fix: switch to an address on the production domain once T1 is done.

T13 (P2). Accessibility basics

There is no skip-to-content link before the nav in `layout()`, and the pre-product banner sits above it. Contrast in `static/style.css` looks sufficient by inspection (white on `#0b5cad` for buttons), but we did not run an automated checker. Fix: add a visually hidden "Skip to content" link targeting `<main id="d6-main">`; run an automated WCAG 2.2 AA check (axe or Lighthouse) once and add the result to the README. Accessibility is a ranking-adjacent quality signal and a legal-risk item, not a ranking factor we can quantify.

T14 (P2). The tests do not guard the new rules

`test.mjs` already checks titles, descriptions, canonicals, one H1, OG and Twitter tags, CSP, scripts, external resources and banned claims. Fix: add checks that draft legal pages carry `noindex` and are absent from the sitemap (T3), that `lastmod` values differ when page dates differ (T4), that `og:image` is present and absolute (T5),

and that no title or JSON-LD name contains "(pre-product)" (T6).

T15 (P3). Performance: nothing to fix, something to measure

The master home page is 8,027 bytes of HTML plus a 2,065-byte stylesheet; the brand build adds two 1,059-byte SVG logos with explicit width and height (no layout shift) and a favicon. There is no JavaScript, no web font and no third-party request; the live home page downloaded 11,261 bytes in 0.20 s from our connection on 2026-10-17 (one request, not a benchmark). Core Web Vitals field data needs real traffic and does not exist yet. Fix: after launch, read Core Web Vitals in Search Console; if the stylesheet ever grows, fingerprint its filename so it can be cached longer than the 3,600 s set in `_headers`.

4. Launch gate

The order matters because some fixes undo others if done early. Secure the domain first (T1), then move to a host that serves the headers (T2), then rebuild with the final `SITE_URL` and the T3 to T6 fixes merged. Only after that should Search Console be verified and the sitemap submitted, followed by a real crawl-based audit of the live domain as `marketing/seo-and-growth.md` section 9 requires. Content from `content-plan.md` starts publishing after the gate, not before, so that no page earns links on a URL we intend to abandon.

- [ ] T1 domain secured; `SITE_URL` set to it; github.io copy redirected or removed
- [ ] T2 headers confirmed with `curl -sI https://<domain>/`
- [ ] T3 to T6 merged; `npm test` passes with the new checks from T14
- [ ] Search Console Domain property verified by DNS; sitemap submitted; coverage report read
- [ ] Baseline crawl audit recorded (indexation, schema validation in the Rich Results Test, internal links)

5. Open questions

- Is scoup.ai registered by the owner through a channel the public RDAP lookup does not show? On 2026-10-17 it did not resolve and the registry RDAP returned 404.
- Was the GitHub Pages deployment from `feat/scoup-brand` intended to be public now? Its README on master says "Not deployed".
- Has the github.io copy already been indexed? We did not check Google's index; Search Console on that URL prefix would answer it.
- Which host will production use (GitHub Pages, Cloudflare Pages, Netlify)? T2's fix depends on it.

Paid and performance channel plan for scoup.ai

> **Status: plan only. The product, scoup.ai, is pre-product.** We have no live website, no ad account history, no traffic, no leads and no paying customers. Every cost figure below is a third-party benchmark with a source line, or a budget we propose, labeled as our decision. No figure in this file is a result we measured. Owner: `digital-marketing` agent, step 1 of the investor package. Last edited 2026-10-17.

Who we sell to first

The owner has decided that scoup.ai sells first to **mid-market shippers** and **SME importers**. Carriers move to a later phase, reached through partners rather than paid media. This plan targets the two first segments only:

- **Mid-market shippers** with roughly $5M to $50M of annual freight spend (an assumed band from the business plan; no public source counts them, per `research/02-industries.md`). Buyers: transportation and logistics managers, AP managers and controllers, and CFOs or VPs of finance, who receive carrier invoices with accessorials on top of the agreed rate. What scoup.ai checks for them: invoice lines against the rate confirmation, unsupported detention and accessorial charges, and fuel surcharge against the carrier's own surcharge table and the diesel index.
- **SME importers**, of which the US Census counts 234,023 (under 500 employees) among 240,535 identified importing companies in 2024; how many import by ocean container is `UNKNOWN — needs primary research`. Buyers: import and logistics managers, customs coordinators, and AP managers or controllers who receive ocean demurrage and detention (D&D) invoices. What scoup.ai checks for them: each D&D invoice against the FMC billing rule, which fixes required invoice content and sets 30-day windows to issue and to dispute (46 CFR part 541; a 2025 court decision vacated only the section on who may be billed).

Source: U.S. Census Bureau, A Profile of U.S. Importing and Exporting Companies, 2023-2024 (CB26-73), https://www.census.gov/foreign-trade/Press-Release/edb/edbrel2024.pdf, accessed 2026-10-17 Source: Thompson Hine, Court Sets Aside Part of FMC's Demurrage and Detention Billing Rule, https://www.thompsonhine.com/insights/court-sets-aside-part-of-fmcs-demurrage-and-detention-billing-rule/, accessed 2026-10-17

The two segments overlap: many mid-market manufacturers and wholesalers both import by container and ship domestically by truck. The financial model's revenue per account is built on exactly that importer-shipper profile (D&D plus fuel surcharge, `financial/assumptions.md`), so an account that buys both checks is the target case, not an edge case. Messaging, capability claims and the name rule follow `creative-concepts.md` and `marketing`'s positioning.

Why paid channels come second

Our go-to-market starts with founder-led design-partner outreach, and `marketing`'s launch order puts no paid acquisition ahead of a partner confirming leakage; paid tests stay locked until milestone G3 (`marketing/marketing-plan.md`). Paid channels here are small, time-boxed message tests, not growth spend. Beyond that lock, four preconditions gate every test: at least one design partner has confirmed leakage on its own closed invoices; the landing page and calculator are live; the conversion events below fire; and at least five design-partner conversations have given us the words buyers use.

One conversion ladder applies to every channel so the numbers stay comparable:

1. **Lead**: a form submission with a work email (design-partner form, calculator, checklist or sample download).

2. **Qualified conversation**: a 20-minute call with someone who approves or disputes freight or D&D invoices.
3. **Leakage review**: the prospect shares closed invoices (20 to 50 loads, or a quarter of D&D invoices) under NDA.
4. **Pilot**: live checks on new invoices, human-approved.
5. **Paying account**: the pilot converts to a paid relationship.

The survival metric is cost per qualified conversation, the earliest step a 6-week test can observe.

Channel 1: LinkedIn Ads (both segments; Year 1)

Audience. Two ad sets. Shipper set: job functions logistics, transportation, supply chain and accounts payable, with titles such as transportation manager, logistics manager, AP manager, controller and CFO or VP of finance, at manufacturers, distributors and retailers in a company-size band that approximates $5M to $50M of freight spend. Importer set: titles such as import manager, logistics manager, customs compliance, AP manager and controller, at manufacturers and wholesalers under 500 employees. LinkedIn exposes neither freight spend nor import volume, so both proxies must be validated on the first leads.

Why it fits. LinkedIn is the one paid platform that targets these job functions directly. It is expensive, and its value shows later in the funnel: Metadata reports $10.20 of pipeline per $1 of LinkedIn lead-gen spend, against $7.97 across all channels in its sample.

Cost benchmarks. CPC $9.39, CPM $63.19, CTR 0.67%, cost per lead $202 (2025 data, all-industry B2B, not freight). The same publisher gives lead-form campaigns at $193 per lead against $346 for landing-page campaigns, document ads at $142 against $265 for video, and a LinkedIn cost per customer of $58,572. The report summary states that the median advertiser paid $376 per lead (cheapest decile $109, most expensive $1,340). Metadata's pages give the sample as 153 advertisers in one place and 138 in another; we treat the figures as directional. We found no published LinkedIn benchmark for freight, logistics or importing.

Source: Metadata.io, What LinkedIn ads cost, https://metadata.io/resources/blog/what-linkedin-ads-cost, accessed 2026-10-17 Source: Metadata.io, 2026 B2B Paid Media Benchmark Report, https://metadata.io/md/benchmark-report-2026.md, accessed 2026-10-17 Source: Webtonic, Trucking and logistics LinkedIn Ads statistics, https://www.webtonic.io/blog/trucking-and-logistics-linkedin-ads-statistics, accessed 2026-10-17

Test budget (our decision). $4,500 over 6 weeks, about $107 per day, split 60/40 between the shipper and importer ad sets ($2,700 and $1,800). Format: document ads (the synthetic accessorial packet, the synthetic D&D invoice check) with LinkedIn lead forms, the two lowest-cost formats in the benchmark. Expected volume at $202: $4,500 / $202 = 22 leads (our arithmetic), too few to estimate conversion rates. The test can kill a message; it cannot prove a channel.

Kill criteria. Stop an ad set if, after half its budget, cost per lead exceeds $376 (the benchmark median), or if its first 10 leads yield no qualified conversation. Stop the channel if, at week 6, cost per qualified conversation exceeds $2,020 (twice the $1,010 implied by $202 / 0.20 in `cac-targets.md`).

Channel 2: Google Search (both segments, high intent only; Year 1)

Audience. People searching buyer-intent phrases. Shipper ad group: accessorial audit, freight invoice audit for mid-size shippers, detention charge dispute, fuel surcharge audit. Importer ad group: demurrage and detention dispute, D&D invoice audit, FMC demurrage invoice requirements, container detention charges dispute. Purely informational queries stay with organic content (`marketing/seo/**`).

Why it fits. Search catches intent and tests headlines fast. The FMC rule gives importer searches a concrete hook (an invoice missing required content, a 30-day dispute window). The weakness is volume; monthly searches for every phrase are `UNKNOWN — needs primary research` (Keyword Planner once an Ads account exists).

Cost benchmarks. LocaliQ's 2026 search benchmarks: Business Services average CPC $5.87, conversion rate 4.85%, cost per lead $93.69; Industrial and Commercial $5.87, 8.20%, $75.19; all industries $5.42, 8.18%, $66.69. These cover a broad, mostly small-business advertiser base. Metadata's B2B sample, counting CRM-matched leads, puts Google Ads at $524 per lead. We plan on the Metadata figure and treat LocaliQ as the optimistic bound.

Source: LocaliQ, 2026 Search Advertising Benchmarks, https://localiq.com/blog/search-advertising-benchmarks/, accessed 2026-10-17 Source: Metadata.io, 2026 B2B Paid Media Benchmark Report, https://metadata.io/md/benchmark-report-2026.md, accessed 2026-10-17 Attempted: WordStream, Google Ads benchmarks, https://www.wordstream.com/blog/ws/2025/04/02/google-ads-benchmarks, returned HTTP 403 on 2026-10-17; not used.

Test budget (our decision). $3,000 over 6 weeks, exact and phrase match, US only, split evenly between the two ad groups, each landing on its solution page. At $5.87 per click the budget buys 511 clicks ($3,000 / $5.87, our arithmetic). If search volume cannot absorb it, the unspent budget is the finding.

Kill criteria. Stop an ad group if cost per lead exceeds $524 after $750, if impressions stay under 500 in its first two weeks, or if no lead becomes a qualified conversation by week 6.

Channel 3: Industry events (Manifest in Year 1; TPM as the importer event)

Audience and fit. For a founder-led company a pre-booked meeting at an event is already a qualified conversation.

- **Manifest** (Las Vegas, February 8 to 10, 2027) targets shippers, logistics technology buyers, startups and investors: the best single fit for the shipper segment and for the seed raise. Partner registration is $1,995 until November 20, 2026, then $2,195 (early), $2,495 (late) and $2,695 on site. Exhibitor prices are not public. A vendor listing for the 2026 edition claimed 7,200+ attendees and 2,000+ shippers; we have not seen a 2027 figure.
- **TPM** (Journal of Commerce, S&P Global; Long Beach, February 28 to March 3, 2027) gathers the ocean container trade, from importers and exporters (beneficial cargo owners) to carriers, forwarders, ports and customs brokers, with "more than 2,300 attendees" per S&P Global. It is the importer segment's natural event. Registration fees are `UNKNOWN — needs primary research` (the TPM site returned HTTP 403 on 2026-10-17).
- **TIA Capital Ideas** (brokers and 3PLs) and **SMC3 JumpStart** (LTL-heavy) fit the later broker and carrier phases, not the first segments, and are not planned.

Source: Manifest, partner page and 2027 media kit, https://partner.manife.st/, accessed 2026-10-17 Source: Exotec, Manifest 2026 event listing, https://www.exotec.com/events/manifest-2026, accessed 2026-10-17 Source: Journal of Commerce, TPM about page (search-result excerpt; site returned 403), https://tpm.joc.com/en/about, accessed 2026-10-17 Source: S&P Global, TPM to celebrate 25th anniversary in 2025 (press release), https://press.spglobal.com/2024-12-16-TPM,-S-P-Globals-Premier-Shipping-and-Supply-Chain-Conference,-to-Celebrate-25th-Anniversary-in-2025, accessed 2026-10-17

Budget (allocation by `marketing`). Year 1: the Manifest pass, $1,995, from marketing programmes; travel (our assumed $1,500) on the separate sales-travel line. TPM 2027 falls in the same month; attending both is a decision for `marketing` once TPM's fee is known. Our recommendation: if the importer segment produces more design-partner conversations than the shipper segment by January 2027, swap Manifest for TPM rather than add a second event.

Kill criteria. Pre-book meetings through the attendee app and outreach; judge the event on qualified conversations held. At $3,495 all-in, Manifest must produce at least 4 ($874 each, our arithmetic: $3,495 / 4), or we skip the next edition.

Channel 4: Trade publications (Year 2)

Audience and fit. JOC reaches the ocean and import trade, the importer segment. DC Velocity reaches distribution and warehouse operators, and FreightWaves reaches shipper logistics staff as well as brokers and carriers. Trade newsletters carry editorial credibility, but they sell at flat rates that resist small tests.

Cost benchmarks. None of these publishers lists rates publicly. FreightWaves calls its options "starting points for conversation" and reports 1.02M unique visitors and 2.36M pageviews a month (search-result excerpt). DC Velocity states it reaches over 330,000 logistics professionals a month. Rates for JOC, DC Velocity and FreightWaves are `UNKNOWN — needs primary research` (request rate cards from S&P Global, Agile Business Media and FreightWaves).

Source: FreightWaves, advertising page (search-result excerpt), https://freightwaves.com/?p=99359, accessed 2026-10-17 Source: Terrapinn, City Freight Show USA partner profile: DC Velocity, https://www.terrapinn.com/conference/city-freight-show-usa/partner-dc-velocity.stm, accessed 2026-10-17

Test budget (our decision, Year 2). A $3,000 cap for one newsletter placement: JOC first if the importer segment leads, DC Velocity if the shipper segment leads. No sponsored articles until we have a measured story.

Kill criteria. Repeat a publication only if its cost per lead is at or below $202 and at least one lead becomes a qualified conversation within 30 days.

Channel 5: Retargeting (both segments; Year 2)

Audience. Visitors to the solution pages, the calculator or the sample packets who did not convert.

Why it fits and cost. Metadata reports retargeting at $196 per lead with $9.60 of pipeline per $1, the best ratio in its sample. The binding constraint is audience size: LinkedIn needs at least 300 matched members before a segment serves, which a new pre-product site may not reach in Year 1. Retargeting is not funded in the Year-1 programmes budget.

Source: Metadata.io, 2026 B2B Paid Media Benchmark Report, https://metadata.io/md/benchmark-report-2026.md, accessed 2026-10-17 Source: LinkedIn, Matched Audiences audience size requirements, https://www.linkedin.com/help/lms/answer/a423690, accessed 2026-10-17

Test budget (our decision, Year 2). $1,000, switched on only when a segment passes 300 matched members. **Kill:** stop if retargeting cost per lead exceeds cold LinkedIn cost per lead in the same period.

Later phase: carriers and partner channels (not budgeted)

Carriers are a later, partner-led segment by owner decision, so every carrier-specific test moves out of Years 1 and 2: participation in carrier Facebook groups and forums, Overdrive and CCJ newsletter slots, and broker-focused events such as TIA Capital Ideas. Meta ads for carriers are dropped entirely for now, per `marketing`. The research behind them stays valid for that phase: 72,053 active interstate for-hire carriers run 5 to 100 power units (FMCSA census queried by `market-researcher`), Metadata reports a Facebook CPM of $15.50 against LinkedIn's $63.19, and detention tools such as FleetOptos and GeoStamp sell through the Geotab and Motive marketplaces, which require an integration scoup.ai does not have.

Partner channels for the first segments are also later: customs brokers and freight forwarders (for importers) and carrier-factoring firms (for carriers). No partner exists today. Per `marketing`'s go-to-market rule, referral terms come only after a pilot and no material implies a partnership before one is signed.

Source: FMCSA, Company Census File (dataset az4n-8mr2), https://data.transportation.gov/resource/az4n-8mr2.json, accessed 2026-10-17 Source: Metadata.io, What LinkedIn ads cost, https://metadata.io/resources/blog/what-linkedin-ads-cost, accessed 2026-10-17

Summary and sequencing

Funding follows the split `finance` set and `marketing` allocated (`marketing/marketing-plan.md`, section 6): Year-1 marketing programmes total $30.0K, with sales travel budgeted separately at $60K.

Channel	Segments	Benchmark CPL used	Test budget	Funding	Window	Primary kill trigger
LinkedIn Ads	Shipper 60%, importer 40%	$202 (median $376)	$4,500	Year 1	6 weeks	CPL > $376 at half budget
Google Search	Shipper and importer ad groups	$524 (LocaliQ $93.69 best case)	$3,000	Year 1	6 weeks	CPL > $524 after $750 per group
Manifest 2027 (attend)	Shipper, investors	n/a	$1,995 pass + $1,500 travel (sales line)	Year 1	one event	< 4 qualified conversations
Trade newsletter (JOC or DC Velocity)	Leading segment	none public	$3,000 cap	Year 2	one send	CPL > $202
Retargeting	Both	$196	$1,000	Year 2	when audience ≥ 300	CPL above cold LinkedIn
Carrier tests (communities, Overdrive/CCJ; Meta dropped)	Carriers	n/a	none	Later phase	n/a	n/a

Year 1, charged to programmes: $4,500 + $3,000 + $1,995 = $9,495, 32% of the $30.0K. Year 2: $3,000 + $1,000 = $4,000, down from $5,500 because the $1,500 carrier Meta test moved to the later phase. TPM 2027 would add an unknown fee if `marketing` chooses it over or alongside Manifest.

Order of operations: (1) outreach and the first five design-partner conversations, in both segments; (2) site, calculator and conversion events live; (3) LinkedIn and Google tests in parallel, two segments each, so we learn which segment responds to paid messages; (4) Manifest (or TPM) in February 2027; (5) in Year 2, one trade placement in the leading segment and retargeting once audiences qualify.

Measurement

Every ad carries UTM tags that include the segment; every form asks the respondent's role and whether they import by container, ship domestically by truck, or both. Qualified conversations are logged in the CRM sheet with source and segment and uploaded back to LinkedIn and Google as offline conversions. Privacy rules follow `marketing/seo-and-growth.md`: EU/EEA consent for every tag, no pre-ticked boxes, and the LinkedIn Insight Tag loads only after consent.

Open questions

- Keyword volumes for the shipper and importer phrases: `UNKNOWN — needs primary research`.
- Number of SME importers that import by ocean container: `UNKNOWN — needs primary research`.
- Count of mid-market shippers by freight spend: `UNKNOWN — needs primary research` (no public dataset).
- TPM 2027 registration fee and whether `marketing` swaps or adds it: `UNKNOWN — needs primary research`.
- Rate cards for JOC, DC Velocity and FreightWaves: `UNKNOWN — needs primary research`.
- A freight- or import-specific paid-media benchmark: none found; our tests will be the first data.

CAC targets per channel for scoup.ai

> **Status: model, not measurement.** The product, scoup.ai, is pre-product with no customers, so no CAC has been observed. Each input is labeled **[sourced]** (third-party benchmark, source line given) or **[assumed]** (our planning assumption, to be replaced by test data). All arithmetic is ours and shown in full. Companion to `channel-plan.md`. Last edited 2026-10-17.

Method

Customer acquisition cost per channel is the channel's cost per lead divided by the share of leads that become paying accounts. We split that share into the four steps of the conversion ladder in `channel-plan.md`, because each step can be measured separately as pilots run, and because the steps fail for different reasons (bad targeting, weak offer, packet quality, pricing).

CAC here is **media CAC**: ad spend or event cost only. It excludes founder time, sales salaries, tooling and the free leakage review we run for every prospect. Fully loaded CAC is higher, and `finance` owns that number.

Funnel assumptions

Step	Rate	Label	Reasoning
Lead to qualified conversation	20%	[assumed]	Many form fills will be students, vendors or companies too small to qualify
Qualified conversation to leakage review	35%	[assumed]	Sharing closed invoices under NDA is a real ask of a stranger
Leakage review to pilot	50%	[assumed]	A prospect who shares files is engaged; the packet must still find money
Pilot to paying account	60%	[assumed]	Contingency pricing lowers the risk of saying yes
Lead to paying account	**2.1 %**	[assumed, product]	0.20 × 0.35 × 0.50 × 0.60 = 0.021

So one paying account needs 1 / 0.021 = 47.6 leads under these assumptions.

A sourced check against our own assumption. Metadata's benchmark reports a LinkedIn cost per lead of $202 and a LinkedIn cost per customer of $58,572 across its B2B sample. Dividing the two gives an observed lead-to-customer rate of $202 / $58,572 = 0.345%, about one sixth of our 2.1%. Their sample is broad B2B software, sales cycles and deal sizes differ from ours, and a contingency offer with a free review may convert better. But the gap is large enough that we carry a pessimistic case built on the 0.345% figure everywhere below. The same report puts the cost of a small-company customer at $31k and a mid-market customer at $103k across all channels.

Source: Metadata.io, What LinkedIn ads cost, https://metadata.io/resources/blog/what-linkedin-ads-cost, accessed 2026-10-17 Source: Metadata.io, 2026 B2B Paid Media Benchmark Report, https://metadata.io/md/benchmark-report-2026.md, accessed 2026-10-17

CAC ceilings: what we can afford

Shipper and importer accounts. By owner decision, scoup.ai sells first to mid-market shippers and SME importers, so these targets apply to those two segments. We align to the current financial model owned by `finance` (`financial/model.xlsx`, Revenue tab, base case; `financial/unit-economics.md`), which replaced the investor brief's earlier $40K / 60% / $25K figures. The model now carries revenue per account by segment **[assumed, model output]**:

Segment (checks bought)	Y1	Y2	Y3	Y5	Y7
SME importer (D&D + fuel surcharge)	$13.0K	$17.9K	$21.9K	$27.6K	$30.8K
Mid-market shipper (fuel surcharge + accessorial and detention audit; 0.77 of the importer amount)	$10.0K	$15.2K	$19.2K	$24.7K	$27.6K
Carrier (later partner channel)	$7.9K	$7.9K	$7.9K	$7.9K	$7.9K

Company gross margin is a model output: 20.6% in Year 2, 60.3% in Year 3, 74.6% in Year 5. A 12-month payback on gross profit allows fully loaded CAC up to revenue per account × gross margin. We reserve two thirds of that for founder and sales time and the free leakage review, so the media CAC target is one third (our decision):

Account	Fully loaded ceiling (revenue × margin)	Media CAC target (÷ 3)	Max CPL (× 0.021)	Max cost per qualified conversation (× 0.105)
Importer, Year 3	$21,900 × 0.603 = $13,206	$4,402	$92	$462
Shipper, Year 3	$19,160 × 0.603 = $11,555	$3,852	$81	$404
Importer, Year 2	$17,900 × 0.206 = $3,687	$1,229	$26	$129
Shipper, Year 2	$15,160 × 0.206 = $3,123	$1,041	$22	$109
Importer buying only the D&D check, Year 3 [assumed]	($10.0K + $85K × 36/52 × 14%) × 0.603 = about $11,000	$3,667	$77	$385
Shipper buying only the fuel check, Year 3 [assumed]	($10.0K + $85K × 0.31 × 14%) × 0.603 = about $8,200	$2,733	$57	$287

The two single-check rows are `finance`'s estimates from the business-case example, not model outputs: a $10.0K Year-3 platform fee plus a share of the $85K Year-3 recovery at 14%, where D&D is $36K of the $52K example (0.69) and fuel surcharge $16K of it (0.31). Shipper figures use the model's unrounded revenue ($19.16K in Year 3, $15.16K in Year 2), as checked by `finance`. The model's plan CAC (all sales and marketing divided by new accounts) is $13.4K in Year 2 and $15.0K in Year 3, base CAC payback on directly sold accounts is 14.5 months in Year 3 and 11.0 in Year 5, and the lifetime value of a directly sold account is $61.9K in Year 3 and $97.5K in Year 5 (`financial/unit-economics.md`).

`finance` also ran our sourced pessimistic case through the model: at a $31K CAC (Metadata's small-company customer cost), LTV:CAC falls to 2.0x in Year 3, payback stretches to 30.1 months, and the plan needs a $9.98M seed and $18.67M of equity in total, against a base seed need of $6.91M (including a 20% buffer) that the planned $7.0M raise covers (`financial/scenarios.md`, `financial/use-of-funds.md`).

Carriers (later phase). Carriers are now a later, partner-led segment, and the model has no separate carrier revenue. The one carrier figure in the package is `market-researcher`'s modeled fee of $7,895 a year for an average 16.7-truck carrier (`research/02-industries.md`; most inputs assumed). At a 60% margin and the same one-third rule, its media CAC target would be $1,579 ($7,895 × 0.60 / 3) and its maximum cost per lead $33 ($1,579 × 0.021), about a sixth of the $202 LinkedIn benchmark. That supports the owner's decision to reach carriers through partners rather than paid media.

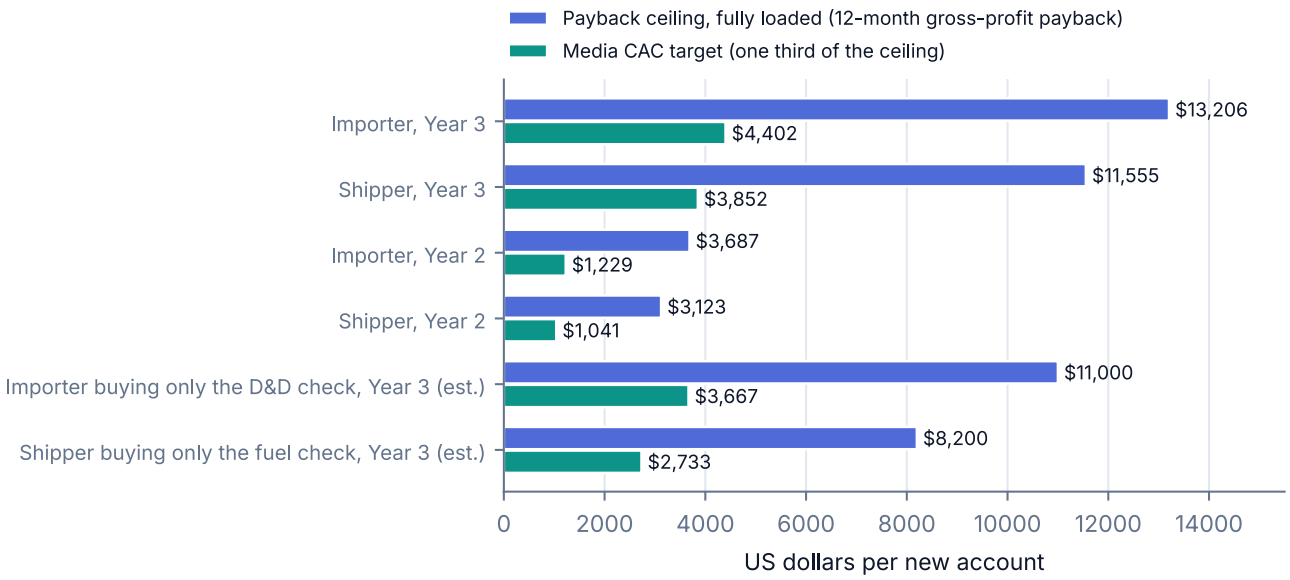


Figure 1: Acquisition-cost target against the payback ceiling for importer and shipper accounts; carriers are not shown because the source gives only a media target for them. Source: marketing/digital/cac-targets.md § CAC ceilings

Modeled CAC per channel

Base case uses our 2.1% lead-to-account rate **[assumed]**. Pessimistic case uses the 0.345% rate implied by Metadata's LinkedIn figures **[sourced, derived]** ($202 / $58,572 = 0.0034487). CPL inputs are cited in `channel-plan.md`.

Channel	CPL input	Label	Base media CAC (CPL / 0.021)	Pessimistic media CAC (CPL / 0.0034487)	Media target Y3 (importer / shipper)	Verdict on paper
LinkedIn Ads, benchmark	$202	[sourced]	$9,619	$58,572	$4,402 / $3,852	About 2.2x the importer and 2.5x the shipper target; needs CPL of $92 / $81
LinkedIn Ads, median advertiser	$376	[sourced]	$17,905	$109,025	$4,402 / $3,852	Over target
Google Search, B2B (Metadata)	$524	[sourced]	$24,952	$151,939	$4,402 / $3,852	Over target
Google Search, LocaliQ Business Services	$93.69	[sourced]	$4,461	$27,166	$4,402 / $3,852	Just above both targets even in the best case
Retargeting	$196	[sourced]	$9,333	$56,832	$4,402 / $3,852	About 2.1x importer and 2.4x shipper target
Manifest 2027 (attend)	$874 per qualified conversation	[assumed: 4 conversations for $3,495]	$8,321	$50,671	$4,402 / $3,852	Over at 4 conversations; at 8 ($3,495 / 8 / 0.105 = $4,161) under the importer target, over the shipper target
Trade newsletter slot	≤ $202 (kill threshold)	[assumed]	≤ $9,619	≤ $58,572	$4,402 / $3,852	Unknown until rate card
Carrier tests (later phase)	not planned	n/a	n/a	n/a	carrier: $1,579	Moved out of Years 1 and 2
Founder outreach and partners	$0 media	[assumed]	$0 media; founder time unpriced	same	n/a	Cheapest on media, capped by founder hours

Event arithmetic: $3,495 / 4 = $873.75 per qualified conversation; the remaining steps convert at 0.35 × 0.50 × 0.60 = 0.105, so $873.75 / 0.105 = $8,321. Pessimistic: conversation-to-account = 0.0034487 / 0.20 = 1.72%, so $873.75 / 0.0172 = $50,671.

Targets we will hold the tests to

At Year-3 economics the targets imply a maximum cost per lead of **$92 for importers and $81 for shippers**, and a maximum cost per qualified conversation of **$462 and $404**. At Year-2 economics, with a 20.6% gross margin, they fall to $26 and $22 per lead. No benchmark channel reaches that, so in Year 2 paid channels remain message tests, and accounts come from outreach, events and referrals.

These are far tighter than the kill criteria in `channel-plan.md` ($376 per lead, $2,020 per qualified conversation), on purpose. Kill criteria decide whether to stop a 6-week test. Targets decide whether a channel deserves budget after the test. A channel can survive its test and still fail its target; it then stays at test-level spend until its conversion steps improve.

The importer account carries the higher target because it buys two checks, D&D and fuel surcharge. A shipper that buys only the fuel check has a media target of $2,733, about 71% of the full shipper figure ($3,852), so paid media should point shippers at the accessorial check, which `marketing` already leads with.

What would change these numbers

Three inputs dominate. Recovered dollars per account drive revenue per account in the financial model, so a smaller recoverable share lowers the CAC we can afford in proportion. Gross margin, a model output that moves from 20.6% in Year 2 to 60.3% in Year 3, triples the affordable CAC between those years. Lead-to-account rate swings CAC sixfold between our base and pessimistic cases, and only real leads can settle it. Check mix matters too: an account that buys one check instead of two can afford about 71% (shipper, fuel only) to 83% (importer, D&D only) of its full-account target, because the platform fee does not shrink with the number of checks. Each test in `channel-plan.md` measures the first funnel step within 6 weeks; the later steps need pilots, which is why we treat every CAC above as a target to test, not an expected result.

Open questions

- Revenue for single-check accounts: `finance` estimates from the business-case example only; not modeled.
- Carrier revenue for the later phase: `UNKNOWN — needs primary research`; `market-researcher` models $7,895 on mostly assumed inputs.
- Funnel step rates: all four [assumed]; first measurement comes from outreach and the paid tests.
- Fully loaded CAC including founder time and the free leakage review: owned by `finance`.
- Recovery-cycle lag between pilot start and first contingency revenue, which stretches payback: owned by `finance`.

Ad and creative concepts for scoup.ai

> **Status: concepts for testing, not approved copy.** The product, scoup.ai, is pre-product: no customers, no measured recovery rate, no measured accuracy. Every concept below passes the project honesty rule by construction, and every one still needs human claims review (and counsel review before launch, per `marketing/seo-and-growth.md`) before it runs. Channels and budgets are in `channel-plan.md`. Last edited 2026-10-17.

Audience and pillars

By owner decision, scoup.ai sells first to **mid-market shippers** (logistics and transportation managers, freight-pay and accounts-payable leads) and **SME importers** (import managers, customs and logistics coordinators). Carrier creative from the earlier draft is withdrawn until the later, partner-led carrier phase.

Each concept serves one of the pillars `marketing` settled (`marketing/brand/README.md`, section 4). **Proof attached** leads every first-touch message: each flagged charge carries its arithmetic and documents. **A person decides** answers the trust objection and runs as secondary creative only, because `market-researcher` found no outside demand for approval controls (`research/04-gaps.md`). **Built for the loads big audit programs skip** frames who we serve without disparaging anyone.

The name appears as lowercase "scoup.ai" and never starts a sentence or headline.

Rules every concept follows

No recovery rate, dollar recovered, accuracy figure or customer line, because we have measured none. Every number is either cited to a primary source or labeled "illustration" with its arithmetic visible. Sample packets and invoices are built on synthetic data and say so on the image. Pre-product status appears in the ad or on the first screen after the click. No competitor is named in paid creative, and no line claims we are the only tool of our kind; BlueCargo and Intelligent Audit already market D&D checks, per `research/01-competitors.md`. Nothing reads as legal advice: D&D copy describes what the FMC rule requires and tells the reader to confirm with counsel or their customs broker.

Two outside facts are approved for use, both re-read today. The FMC reports that the nine largest ocean carriers "collected roughly $15.4 billion in D&D charges between April 1, 2020, and March 31, 2025"; we present that as charges billed, never as money anyone can recover. Thompson Hine summarizes the FMC rule this way: "Failure of the billing party to abide by the invoice content and timing requirements relieves the billed party's obligation to pay an invoice for detention and demurrage charges", and the billed party has 30 days to dispute.

Source: Federal Maritime Commission, Detention and Demurrage, https://www.fmc.gov/detention-and-demurrage/, accessed 2026-10-17 Source: Thompson Hine, FMC Issues Final Rule on Billing Requirements and Practices for Detention and Demurrage Invoices, https://www.thompsonhine.com/insights/fmc-issues-final-rule-on-billing-requirements-and-practices-for-detention-and-demurrage-invoices/, accessed 2026-10-17

Capability claims are limited to what the product does today on synthetic data: the D&D audit checks invoice timing (46 CFR 541.7) and required content (46 CFR 541.6) field by field; the fuel check recomputes a surcharge from the carrier's own table and the active diesel index; the accessorial check compares invoice lines with the rate confirmation and holds uncertain lines for review (`freight-recovery-build`, `packages/shared/src/ocean-dd`, `api/src/intelligence/fuel.ts`).

Concept 1: "Off AP's desk" (lead test)

- **Pillar:** Proof attached. **Audience:** shippers; AP manager and controller titles first, then transportation/logistics managers and CFOs/VPs finance. **Channels:** LinkedIn document ad, Google Search.
- **Headline:** Take accessorial review off AP's desk. Every line checked against the rate confirmation, with the source attached.
- **Alternate headline (A/B):** Stop paying accessorials that were never authorized.
- **Why workload leads:** RXO's November 2024 survey of 750 shippers and carriers (`research/04-gaps.md`, gap 5) found most shippers call accessorial requests a top accounting time-waster while most say carriers are usually honest, which favors a workload message over an overbilling one. Any creative that quotes RXO carries the source line below and presents the figures as RXO's survey findings, never as scoup.ai results.
- **Body:** Each invoice line is checked against the rate confirmation. Charges the paperwork supports pass; charges it does not support come back with the page that shows why. Uncertain lines are held for your review, not disputed.
- **Visual:** a synthetic invoice with one line highlighted and a margin note "not on rate con: review", stamped SYNTHETIC.
- **CTA:** See a sample accessorial check.
- **Search ad variant:** "Accessorial audit, with the proof | Every line checked against the rate con | Pre-product: design partners wanted".
- **Must not say:** "cut freight costs by X%", any savings figure, any claim about how often charges are unauthorized, or RXO's percentages without attribution.

Source: RXO, What Shippers & Carriers Really Think About Trucking Accessorials (Original Research), https://rxo.com/resources/research/freight-accessorials/, accessed 2026-10-17

Concept 2: "The 30-day window" (lead test)

- **Pillar:** Proof attached. **Audience:** SME ocean importers (import/logistics managers, AP managers and controllers). **Channels:** LinkedIn document ad, Google Search, JOC (Year 2).
- **Headline:** Check every D&D invoice before the 30-day window closes.
- **Alternate headline (A/B):** A D&D invoice has rules. So does the clock to dispute it.
- **Body:** Under the FMC billing rule, a demurrage or detention invoice must carry required information and arrive on time, and you have 30 days to dispute it. scoup.ai's check (in design-partner validation) reads each invoice field by field, cites the section it fails, and tracks your dispute deadline. Confirm any dispute with your counsel or customs broker.
- **Visual:** a synthetic D&D invoice with two fields flagged ("container number missing: 46 CFR 541.6") and a 30-day countdown bar.
- **CTA:** See a sample D&D invoice check.
- **Citation rule:** any line stating the rule's effect (that a non-compliant invoice need not be paid) carries the Thompson Hine citation; the gap and its sources are in `research/04-gaps.md`, gap 3.
- **Search ad variant:** "D&D invoice check, by the FMC rule | Required fields and 30-day deadlines tracked | Pre-product pilot".
- **Must not say:** "you don't have to pay", "get your money back", "guaranteed refund".

Concept 3: "No source, no dispute" (lead test)

- **Pillar:** Proof attached. **Audience:** both segments. **Channels:** LinkedIn single-image and document ads.
- **Headline:** No source, no dispute.
- **Alternate headline (A/B):** Every flagged line points to a page.
- **Wording rule:** keep "claim" out of headlines and ad titles; `seo` found buyers read it as cargo loss and damage.

- **Body:** If a charge cannot be traced to a rate confirmation, a tariff table or a rule section, it does not go in the dispute packet. Software reads the documents, rules do the arithmetic, a person approves the send.
- **Visual:** one packet line ("Fuel surcharge billed $412; table gives $371", labeled illustration) with a callout to the synthetic surcharge table row and diesel index week.
- **CTA:** See a sample packet (synthetic data).
- **Must not say:** "accurate", "error-free", "AI-powered".

Concept 4: "Fuel surcharge, recomputed"

- **Pillar:** Proof attached. **Audience:** shippers. **Channels:** LinkedIn (second wave), Google Search.
- **Headline:** Your carrier's fuel table, applied to the week you shipped.
- **Body:** The fuel check takes the surcharge schedule you agreed, the ship date and the published diesel price, and recomputes the charge line by line. Differences come with the arithmetic.
- **Visual:** a three-row worked example labeled "illustration": index price, table bracket, surcharge billed versus recomputed.
- **CTA:** Run a sample fuel check.
- **Must not say:** "carriers overcharge you", any overbilling percentage.

Concept 5: "Too small for an audit program, too big to ignore"

- **Pillar:** Built for the loads big audit programs skip. **Audience:** shippers with $5M to $50M of freight spend. **Channels:** LinkedIn, Manifest follow-up email.
- **Headline:** A freight audit sized for a $20M freight budget.
- **Body:** No integration project and no enterprise contract. Send last quarter's carrier invoices and rate confirmations; get back the lines the paperwork does not support, each with its proof. Pre-product: we are looking for design partners.
- **Visual:** a short stack of invoices going in, one packet coming out.
- **CTA:** Apply for a free leakage review.
- **Honesty check:** says who we build for, not how others treat them. "No integration project" is true because no integration exists; the FAQ says so. $20M is the business plan's illustrative midpoint, not a customer.
- **Must not say:** "the big audit firms ignore you", "integrates with your TMS".

Concept 6: "Import by container, ship by truck?"

- **Pillar:** Built for the loads big audit programs skip. **Audience:** manufacturers and wholesalers that both import and ship domestically. **Channels:** LinkedIn importer ad set, TPM or Manifest follow-up.
- **Headline:** Ocean D&D and truck accessorials, checked in one review.
- **Body:** One design-partner review covers both sides of your freight: D&D invoices against the FMC rule, carrier invoices against your rate confirmations and fuel table. Under NDA, read-only, on invoices you have already closed.
- **Visual:** a container and a trailer feeding the same packet.
- **CTA:** Apply for a free leakage review.
- **Honesty check:** describes the importer-shipper account the financial model assumes; both checks exist on synthetic data.
- **Must not say:** "one platform for all your freight", any combined savings figure.

Concept 7: "We're pre-product. Test us on your closed invoices." (founder content)

- **Pillar:** Proof attached. **Audience:** both. **Channels:** founder LinkedIn posts from day one; lead-form ad if posts produce conversations.

- **Headline:** Send us a quarter of closed invoices. We'll show you what the checks find, or that they find nothing.
- **Body:** We are in design-partner validation. Under NDA, read-only, we run the checks on invoices you have already paid and walk you through every flagged line. No software to buy.
- **Visual:** founder photo (with consent) and a plain text card.
- **Must not say:** "limited spots", countdowns, any number of partners we do not have.

Concept 8: "The packet, unboxed"

- **Pillar:** Proof attached. **Audience:** both, with a shipper and an importer version. **Channels:** LinkedIn document ad, gated download.
- **Headline:** What a dispute looks like with the proof attached.
- **Body:** Page through a complete sample: the invoice, the rule or rate-con term each flagged line breaks, the arithmetic, and the draft dispute letter. Built on synthetic data so you can see everything.
- **Visual:** the first three pages, each stamped "SYNTHETIC SAMPLE".
- **Must not say:** "real dispute", "won", any outcome.

Secondary creative

Concept 9, "You approve every dispute" (A person decides): "Nothing is sent without your OK. You read the draft, edit it and decide." It runs on the solution pages and in follow-up emails as the answer to "will this damage my carrier or ocean-carrier relationship?", never as a lead ad.

Concept 10, "You ran the numbers" (Proof attached, retargeting): "You estimated it from your inputs. Want it run on your invoices?" It is the retargeting ad for calculator visitors and waits for Year 2, when retargeting is funded.

How we test them

Year 1 LinkedIn ($4,500): the shipper ad set runs Concepts 1 and 3 ($1,350 each of its $2,700), the importer ad set runs Concepts 2 and 6 ($900 each of its $1,800). Google Search ($3,000) uses the search variants of Concepts 1 and 2, one per ad group. Concept 7 runs as founder content from day one at no media cost. Concepts 4, 5 and 8 enter the second wave or replace any lead concept whose first 10 leads produce no qualified conversation. Concept 9 lives on pages and emails; Concept 10 waits for Year-2 retargeting. We compare concepts on cost per qualified conversation, not click-through rate.

Open questions

- Counsel review of the FMC-rule copy in Concepts 2 and 6, so it describes the rule without reading as legal advice.
- Founder photo and named-person consent for Concept 7.
- Whether the importer version of Concept 8 needs a customs-broker reviewer before publication.